

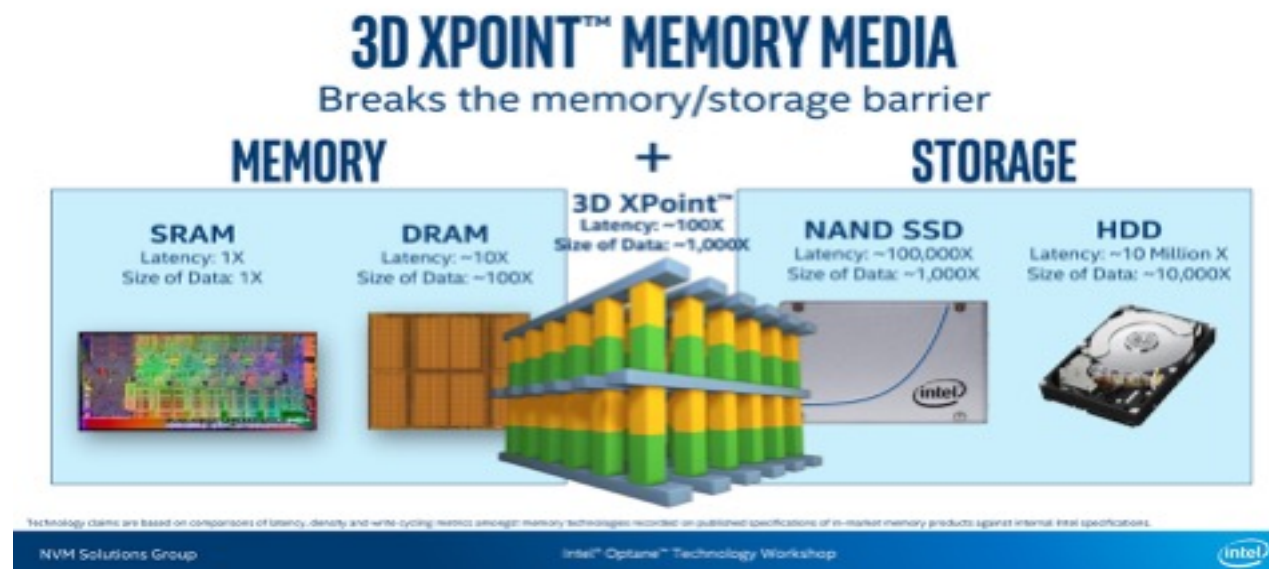
# Distributed Shared Persistent Memory (SoCC '17)

*Yizhou Shan, Yiyang Zhang*


 *WukLab*



# Persistent Memory (PM/NVM)



Byte Addressable  
Persistent  
Low Latency  
Capacity  
Cost effective



# Many PM Work, but All in Single Machine

- Local memory models
  - NV-Heaps [ASPLOS '11], Mnemosyne [ASPLOS '11]
  - Memory Persistency [ISCA '14], Synchronous Ordering [Micro'16]
- Local file systems
  - BPFS [SOSP'09], PMFS [EuroSys'14], SCMFS [SC'11], HiNFS [EuroSys'16]
- Local transaction/logging systems
  - NVWAL [ASPLOS'16], SCT/DCT [ASPLOS'16], Kamino-Tx [Eurosys'17]

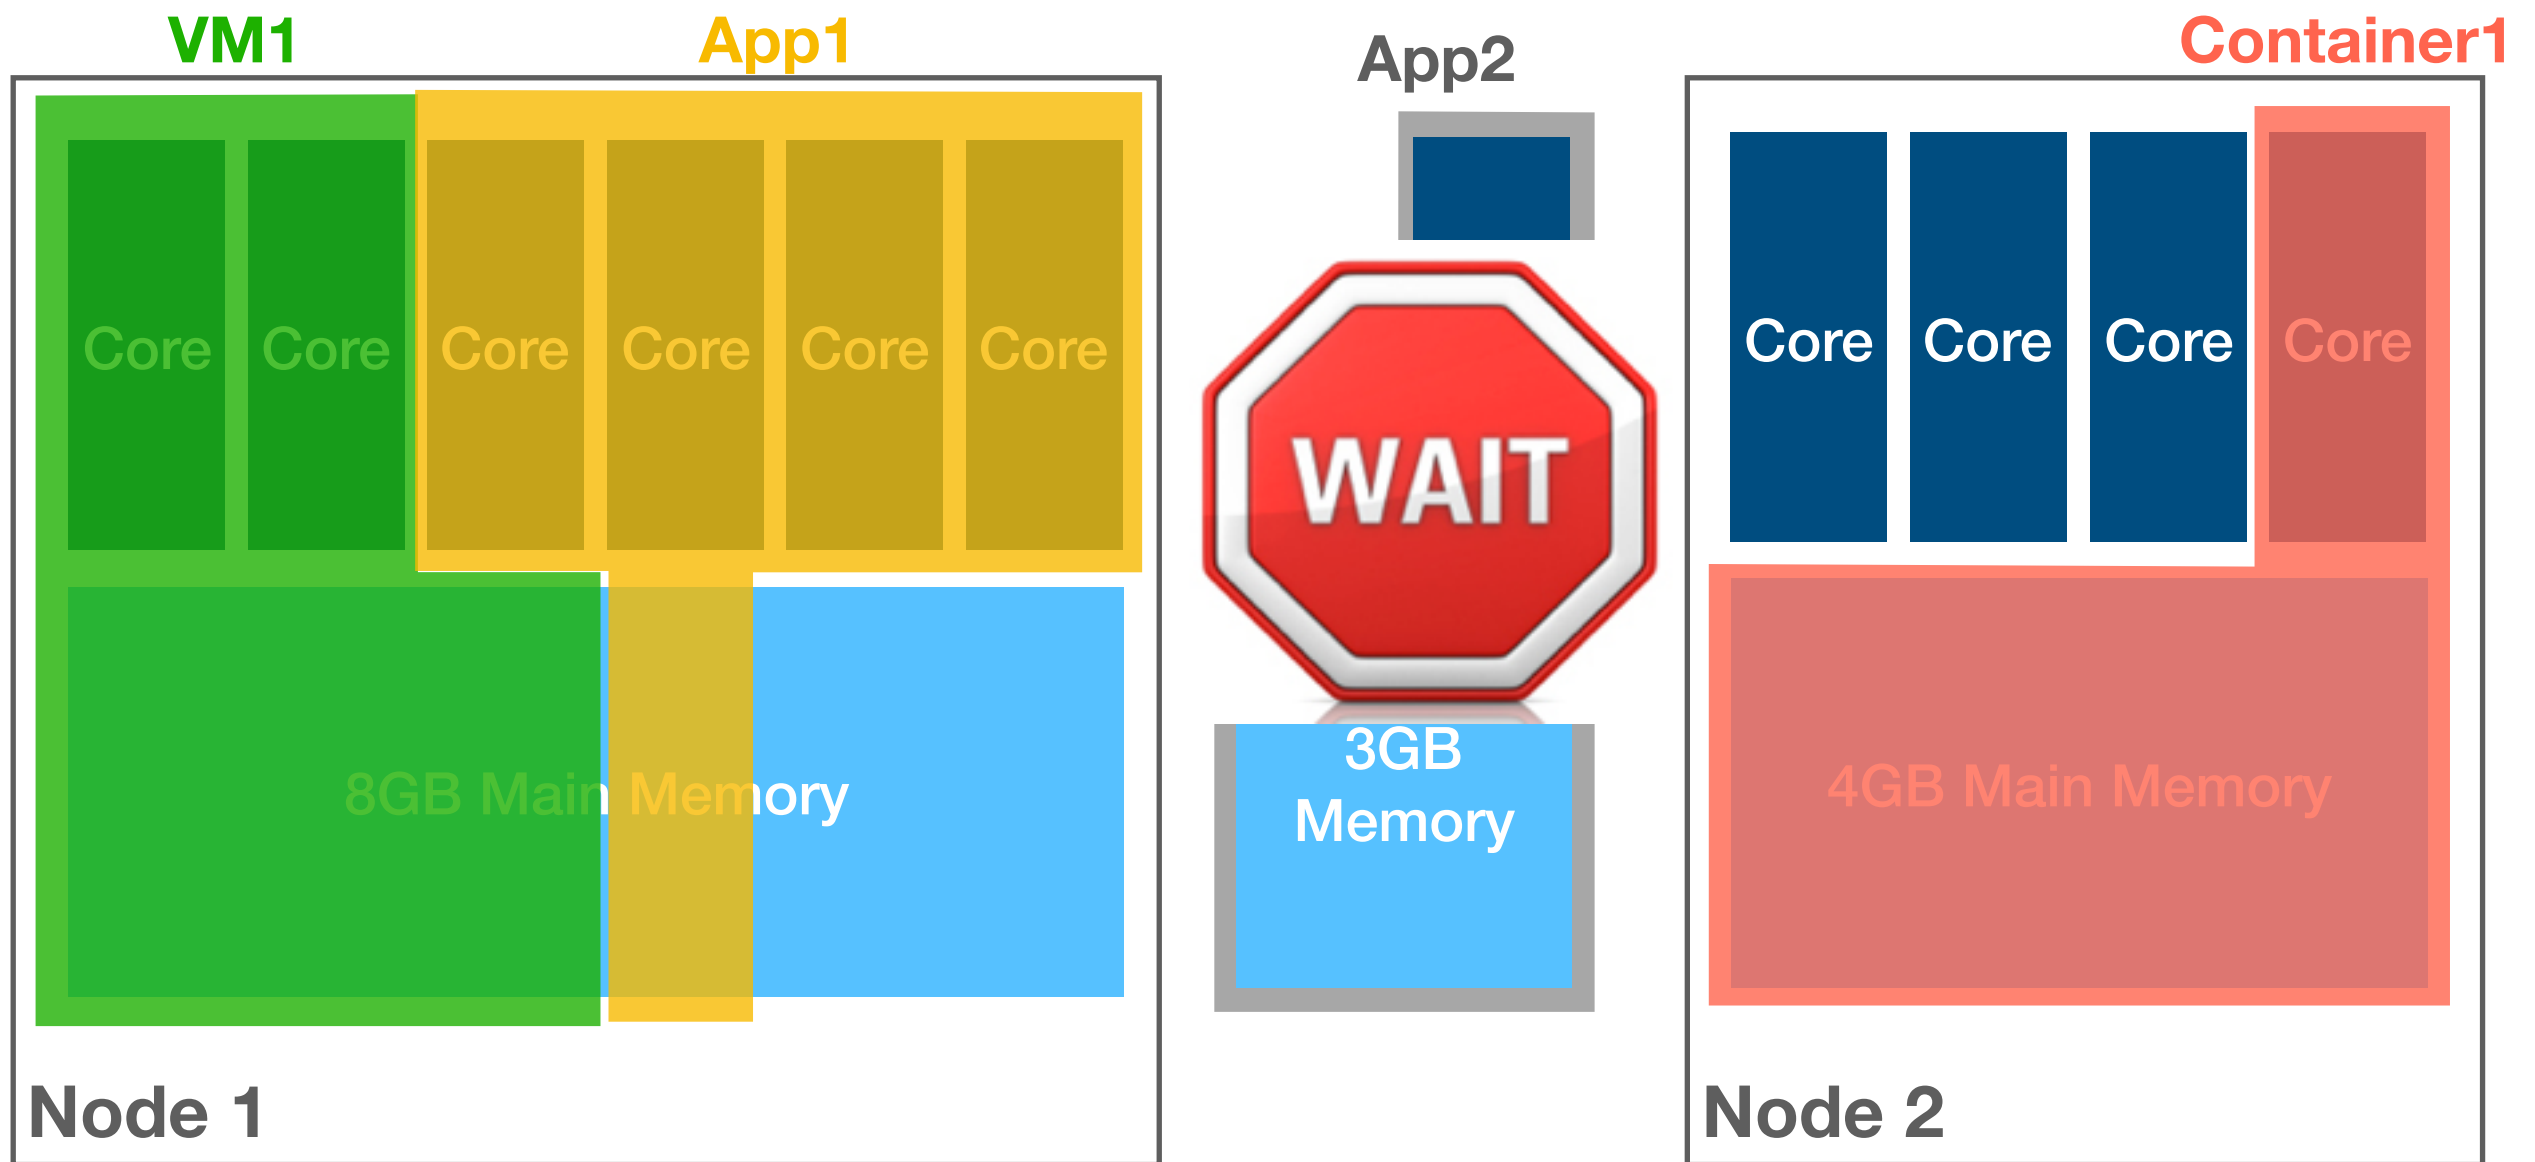
# Moving PM into Datacenters

- PM fits datacenter
  - Applications require a lot memory
  - and accessing persistent data fast
  - with low monetary cost
- Challenges
  - Handle node failure
  - Ensure good performance and scalability
  - Easy-to-use abstraction


# How to Use PM in Distributed Environments?

- As distributed memory?
- As distributed storage?
- Mojim [*Zhang et al., ASPLOS'15*]
  - **First** PM work in distributed environments
  - Efficient **PM replication**
  - *But far from a full-fledged distributed NVM system*

# Resource Allocation in Datacenters



# Resource Utilization in Production Clusters



\* Google Production Cluster Trace Data. "<https://github.com/google/cluster-data>"


**Unused Resource + Waiting/Killed Jobs Because of Physical-Node Constraints**

**Q1: How to achieve better  
resource utilization?**


**Use remote memory**



# Distributed (Remote) Memory




# Modern Datacenter Applications Have Significant Memory Sharing



PowerGraph

TensorFlow



**Q2: How to scale out  
parallel applications?**

**Distributed shared  
memory**

# What about persistence?

- Data persistence is useful
  - Many existing data storage systems
    - ➔ Performance
  - Memory-based, long-running applications
    - ➔ Checkpointing

**Q3: How to provide data  
persistence?**

# DSM

**Distributed Shared Persistent  
Memory (DSPM)**

**a significant step towards  
using PM in datacenters**

# DSPM

- Native memory load/store interface
  - Local or remote (transparent)
  - Pointers and in-memory data structures
- Supports memory read/write sharing

# DSM

**Distributed Shared Persistent  
Memory (DSPM)**

**a significant step towards  
using PM in datacenters**



# DSPM

## DSPM: One Layer Approach

Benefits of both memory and storage  
No redundant layers


- No data marshaling/unmarshalling

# *Hotpot:* A Kernel-Level RDMA-Based DSPM System

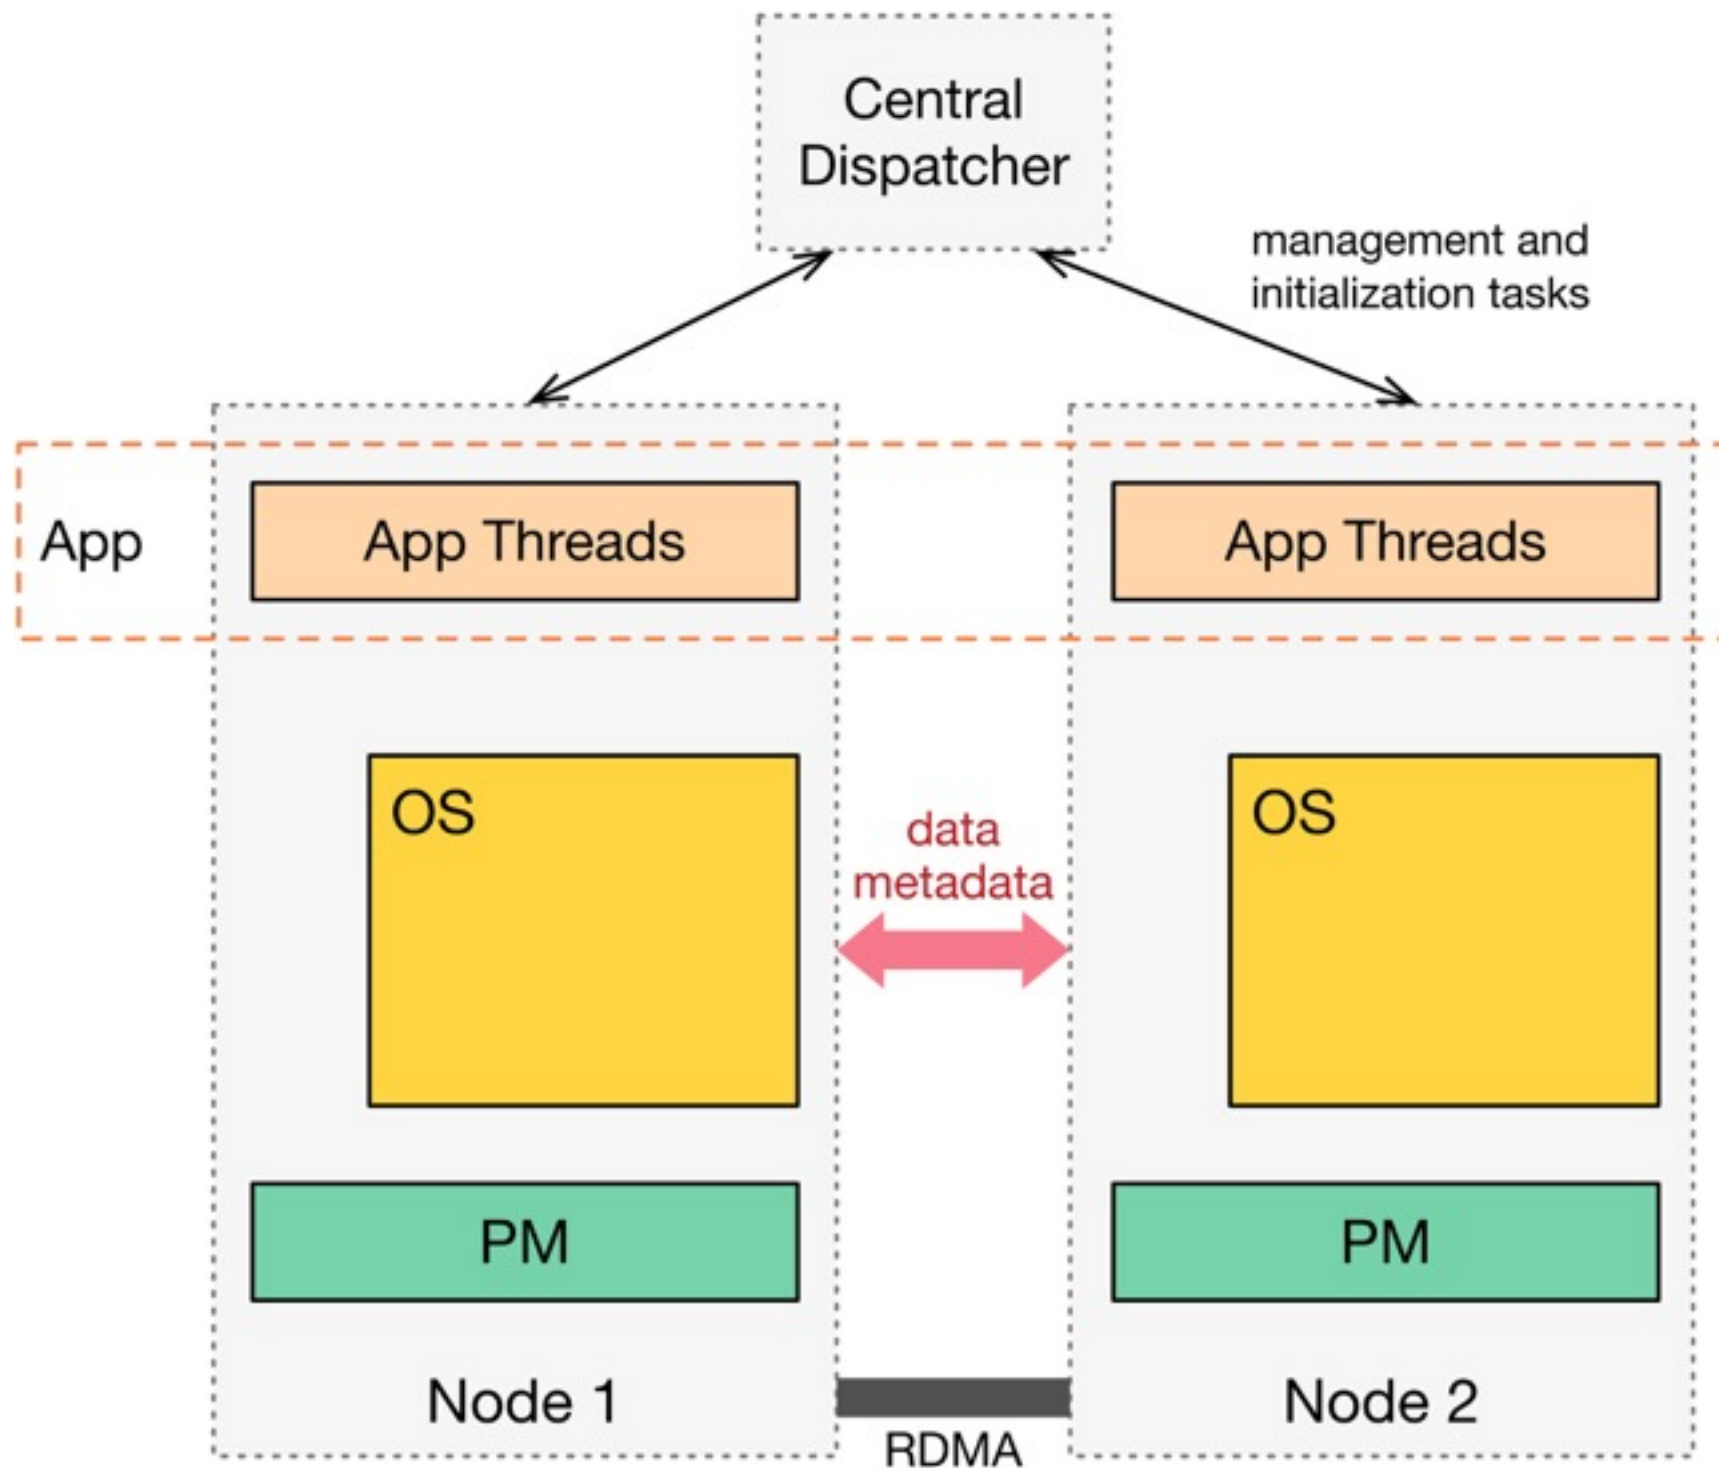
- Easy to use
- Native memory interface
- Fast, scalable
- Flexible consistency levels
- Data durability & reliability



# Hotpot Architecture

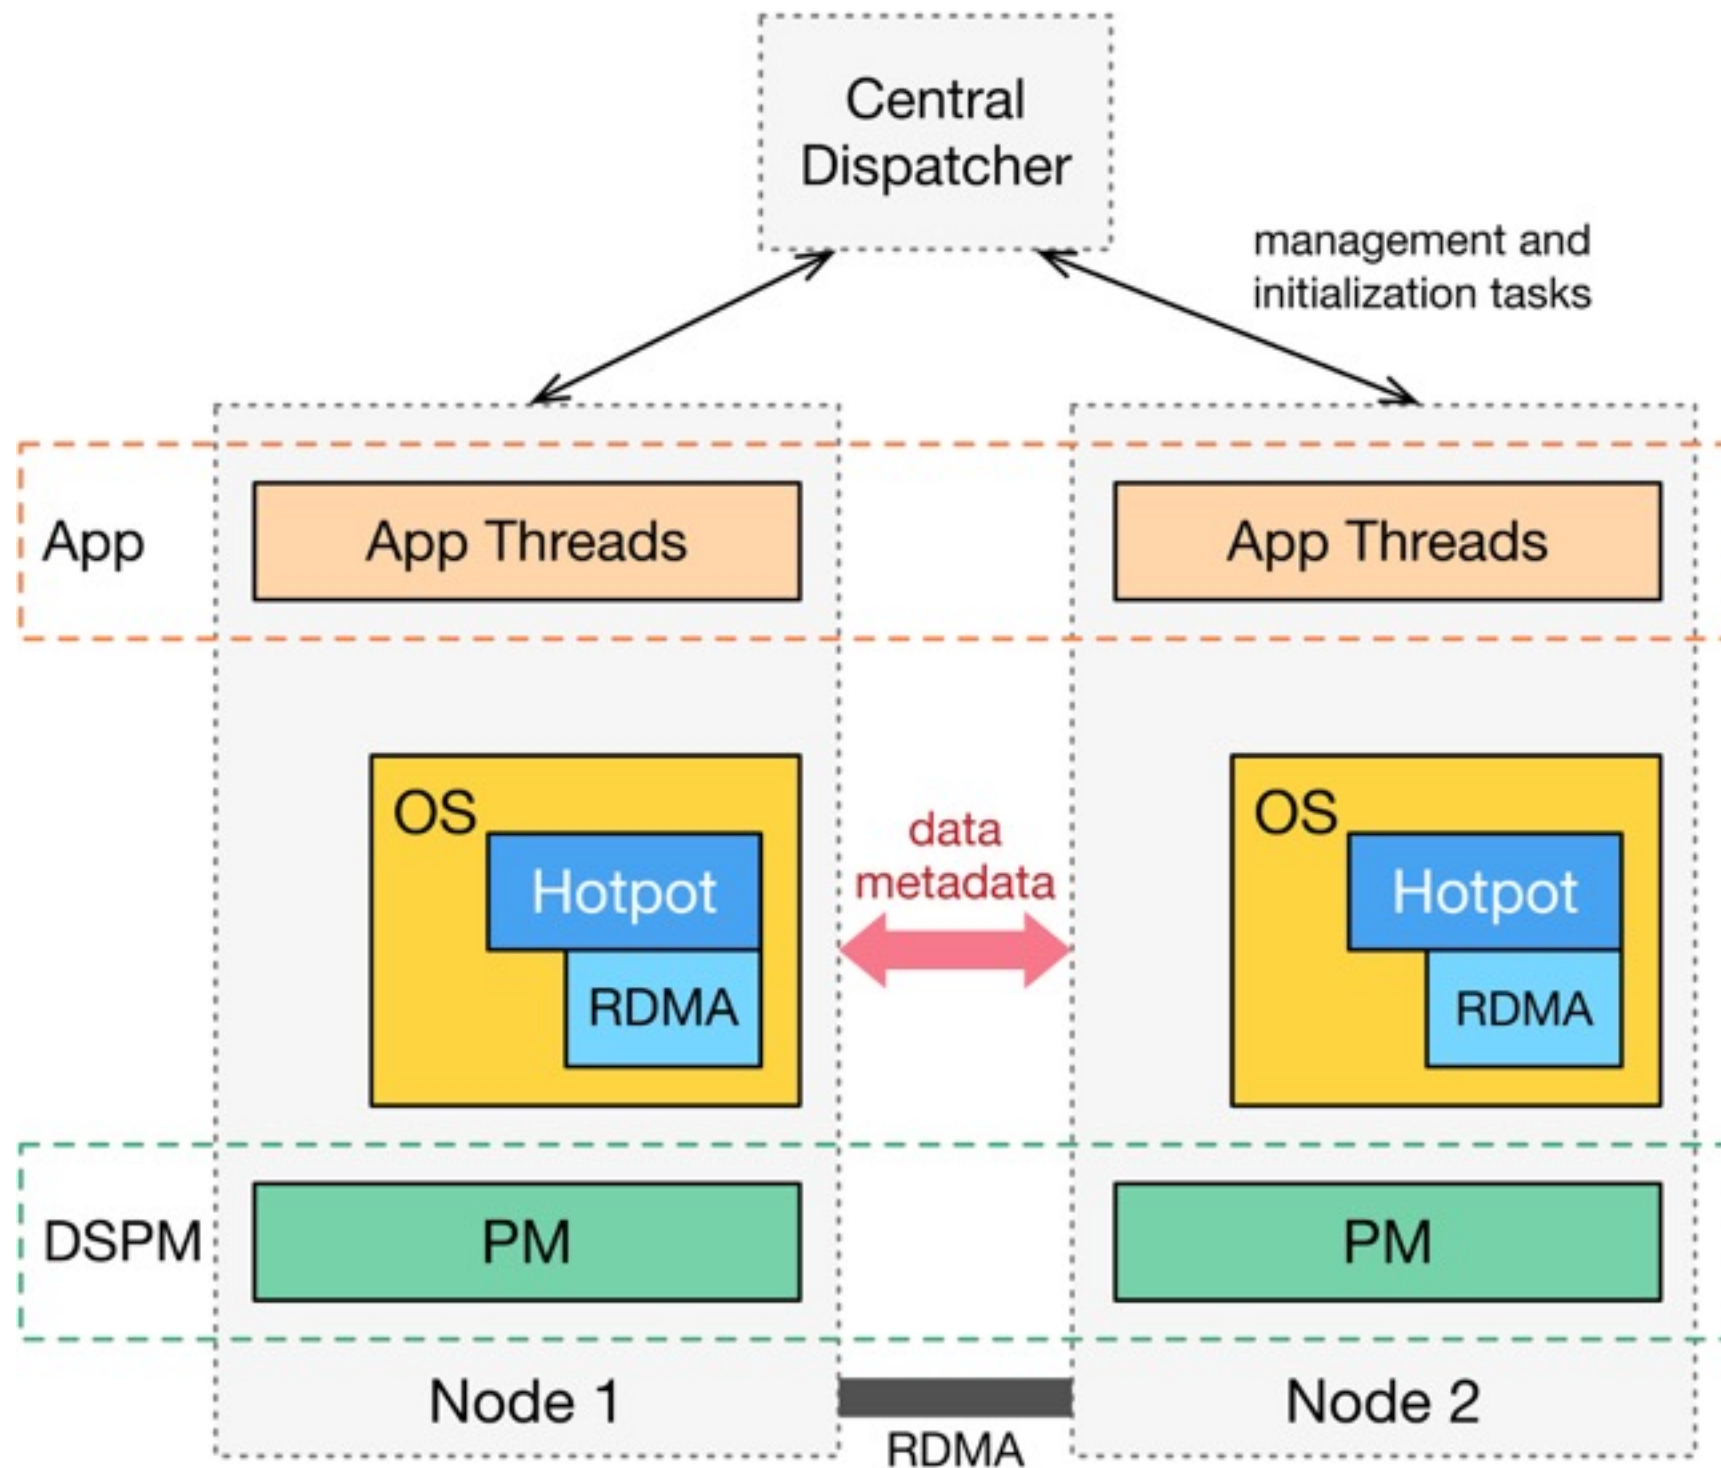


# Hotpot Architecture

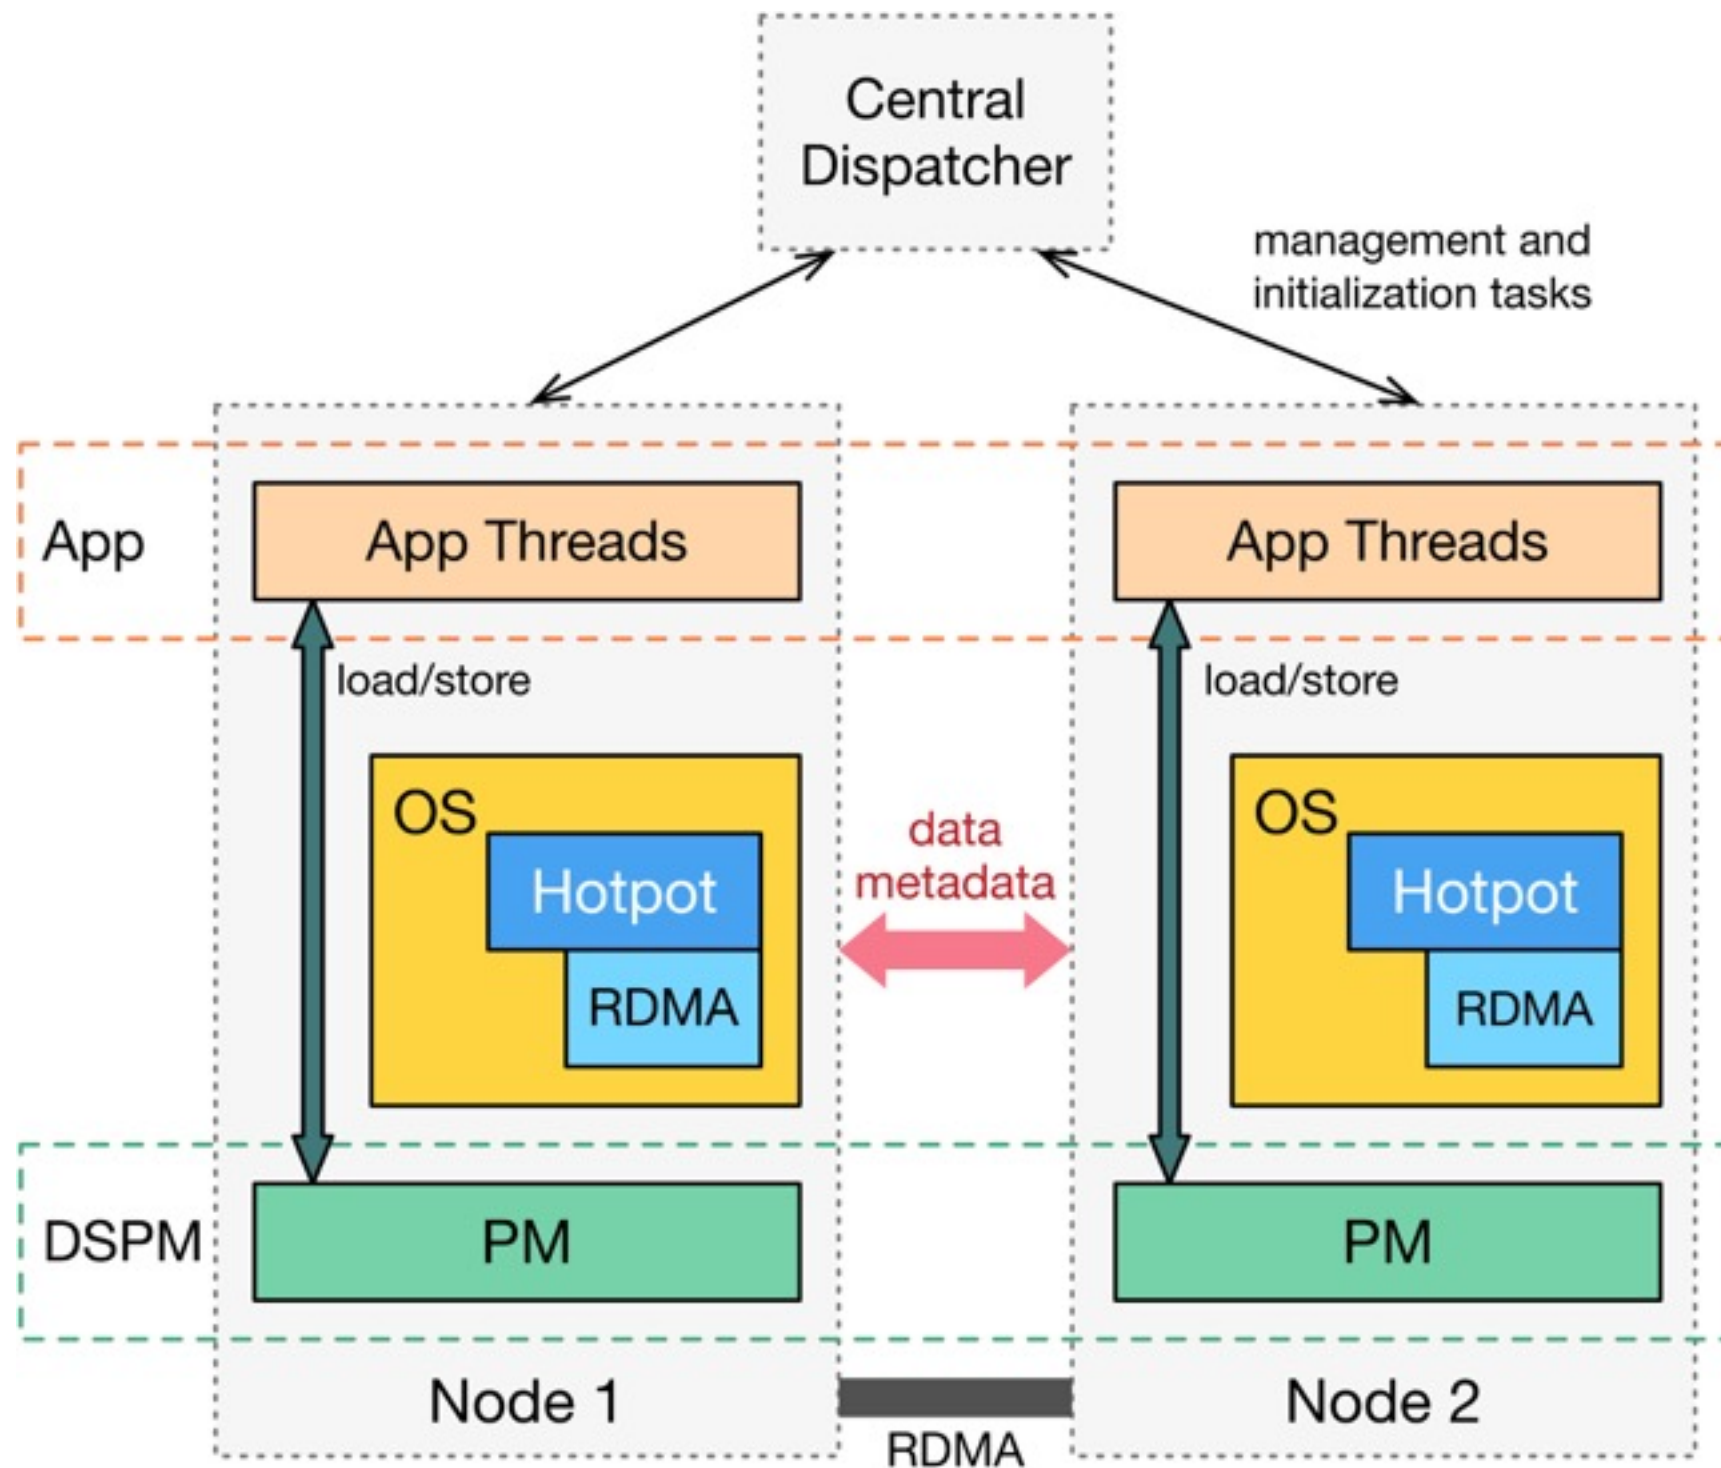




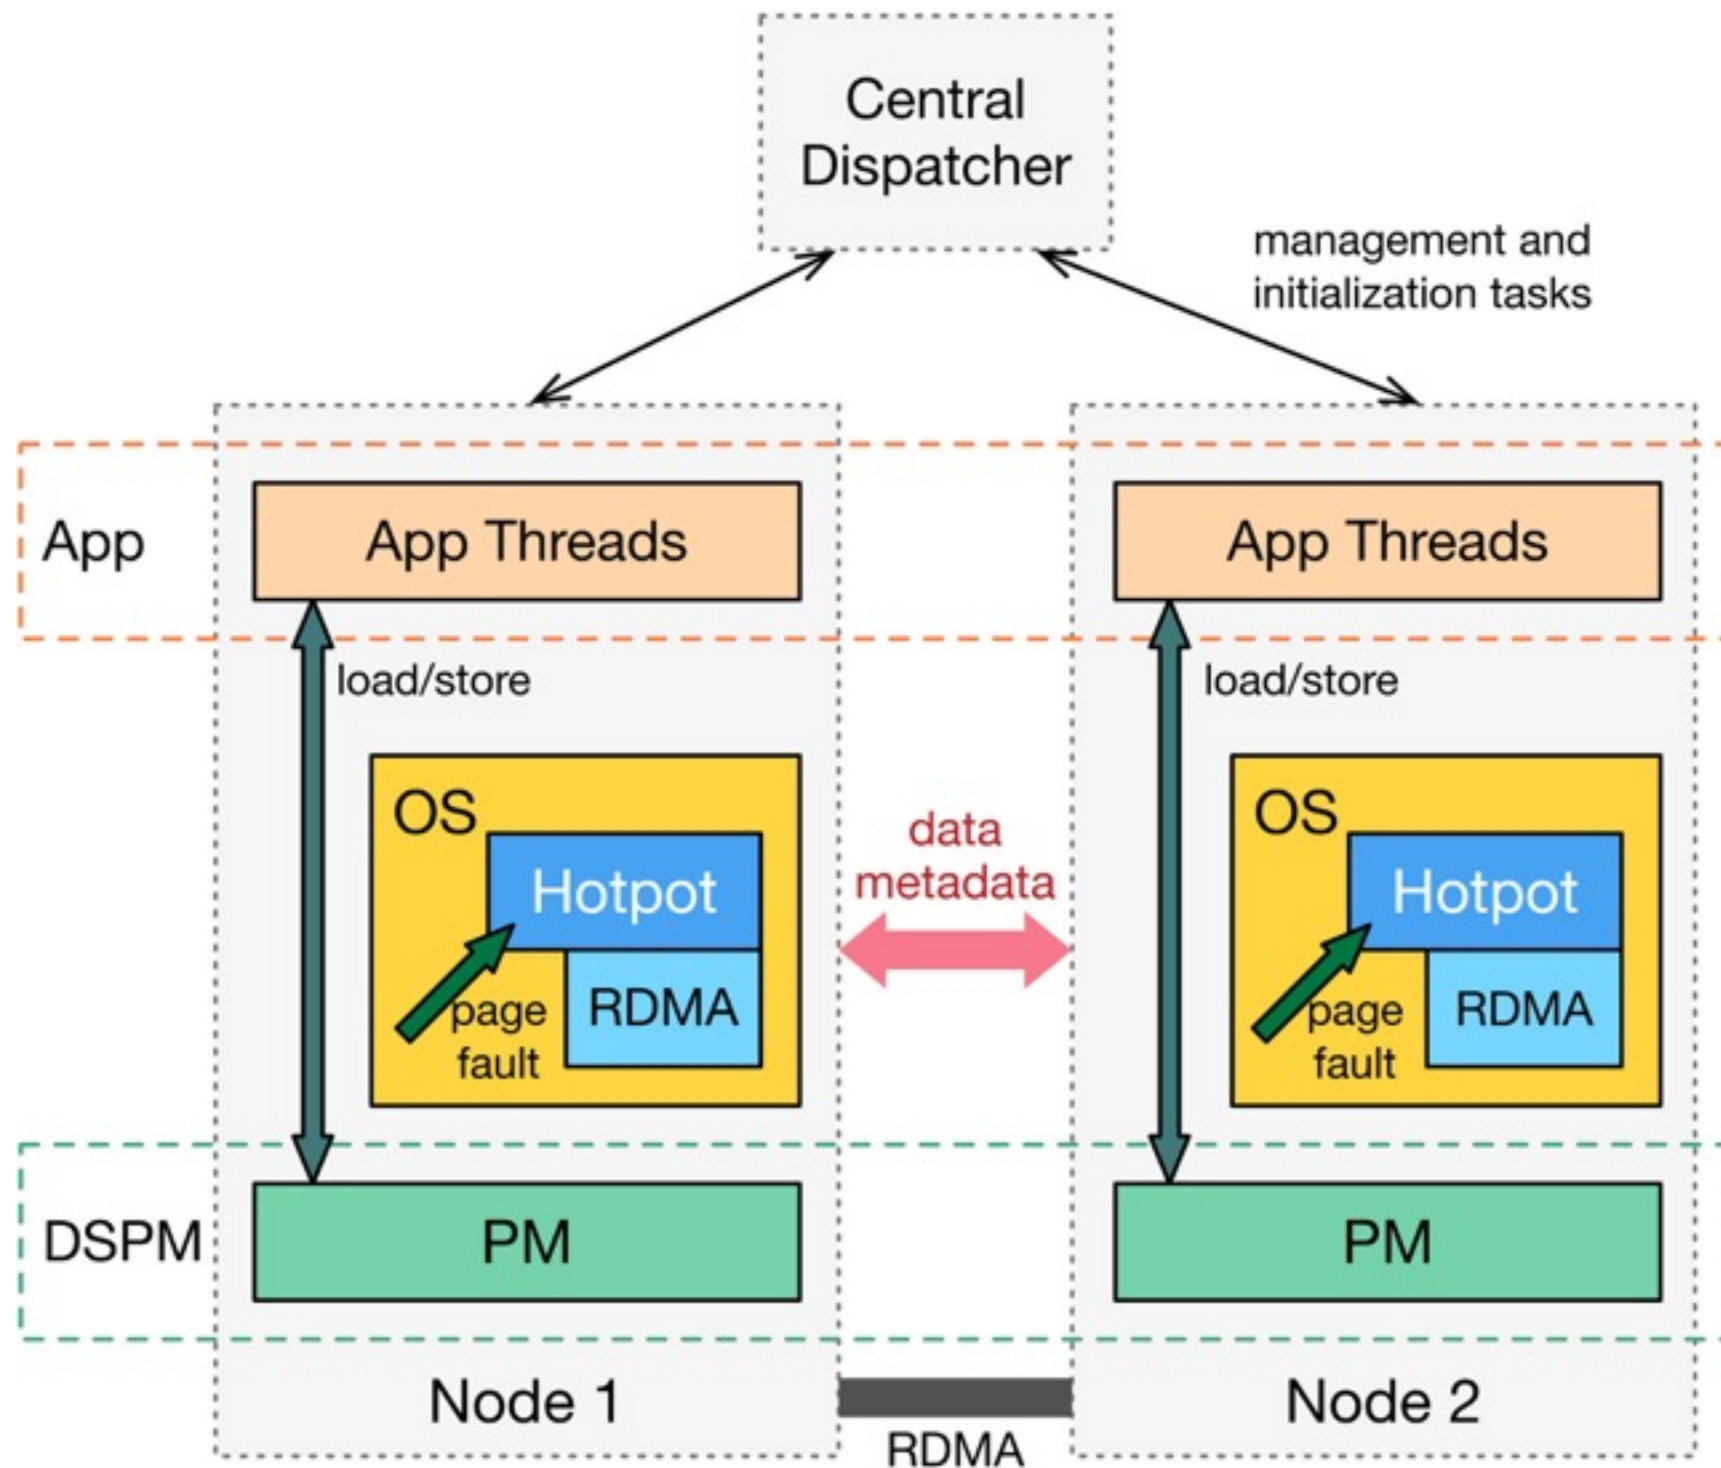
# Hotpot Architecture



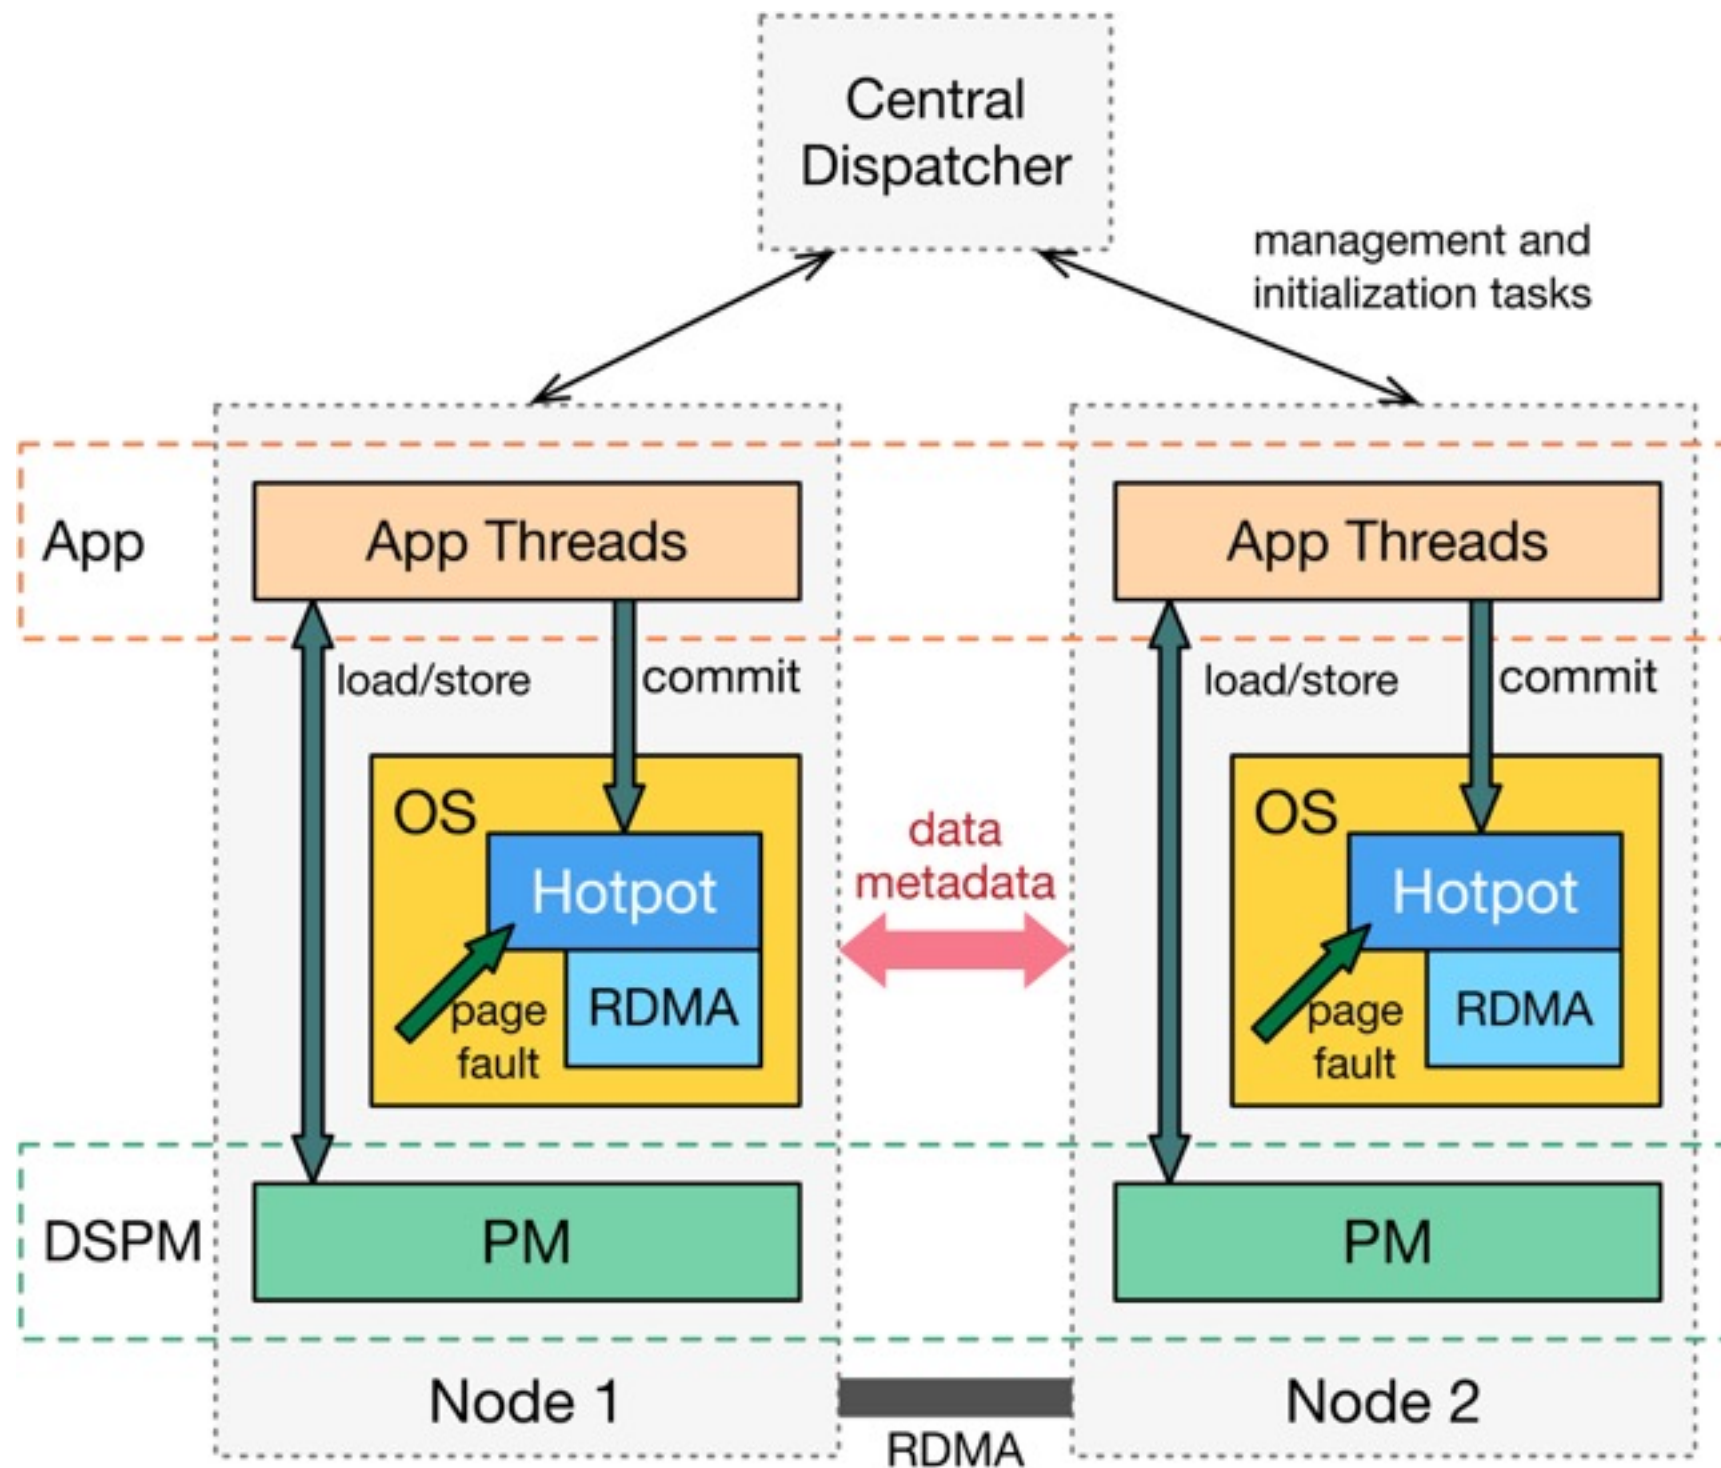
# Hotpot Architecture



# Hotpot Architecture



# Hotpot Architecture





# Hotpot Code Example

```
/* Open a dataset named 'boilermaker' */  
int fd = open("/mnt/hotpot/boilermaker", O_CREAT|O_RDWR);  
/* map it to application's virtual address space */  
void *base = mmap(0, 40960, PROT_WRITE, MAP_PRIVATE, fd, 0);  
  
/* First access: Hotpot will fetch page from remote */  
*base = 9;  
  
/* Later accesses: Direct memory load/store */  
memset(base, 0x27, PAGE_SIZE);  
  
/* Commit data: making data coherent, durable, and replicated */  
msync(sg_addr, sg_len, MSYNC_HOTPOT);
```

# How to efficiently add P to “DSM”?

- Distributed Shared Memory
  - Cache remote memory **on-demand** for fast local access
  - Multiple **redundant** copies
- Distributed Storage Systems
  - **Actively** add more **redundancy** to provide data reliability




**Integrate two forms of redundancy  
with *morphable page states***


**One Layer Principle**

# Morphable Page States

- A PM page can serve different purposes, possibly at different times
  - as a local cached copy to improve performance
  - as a redundant data page to improve data reliability



*Node 2 accesses page 3*



# How to efficiently add P to “DSM”?

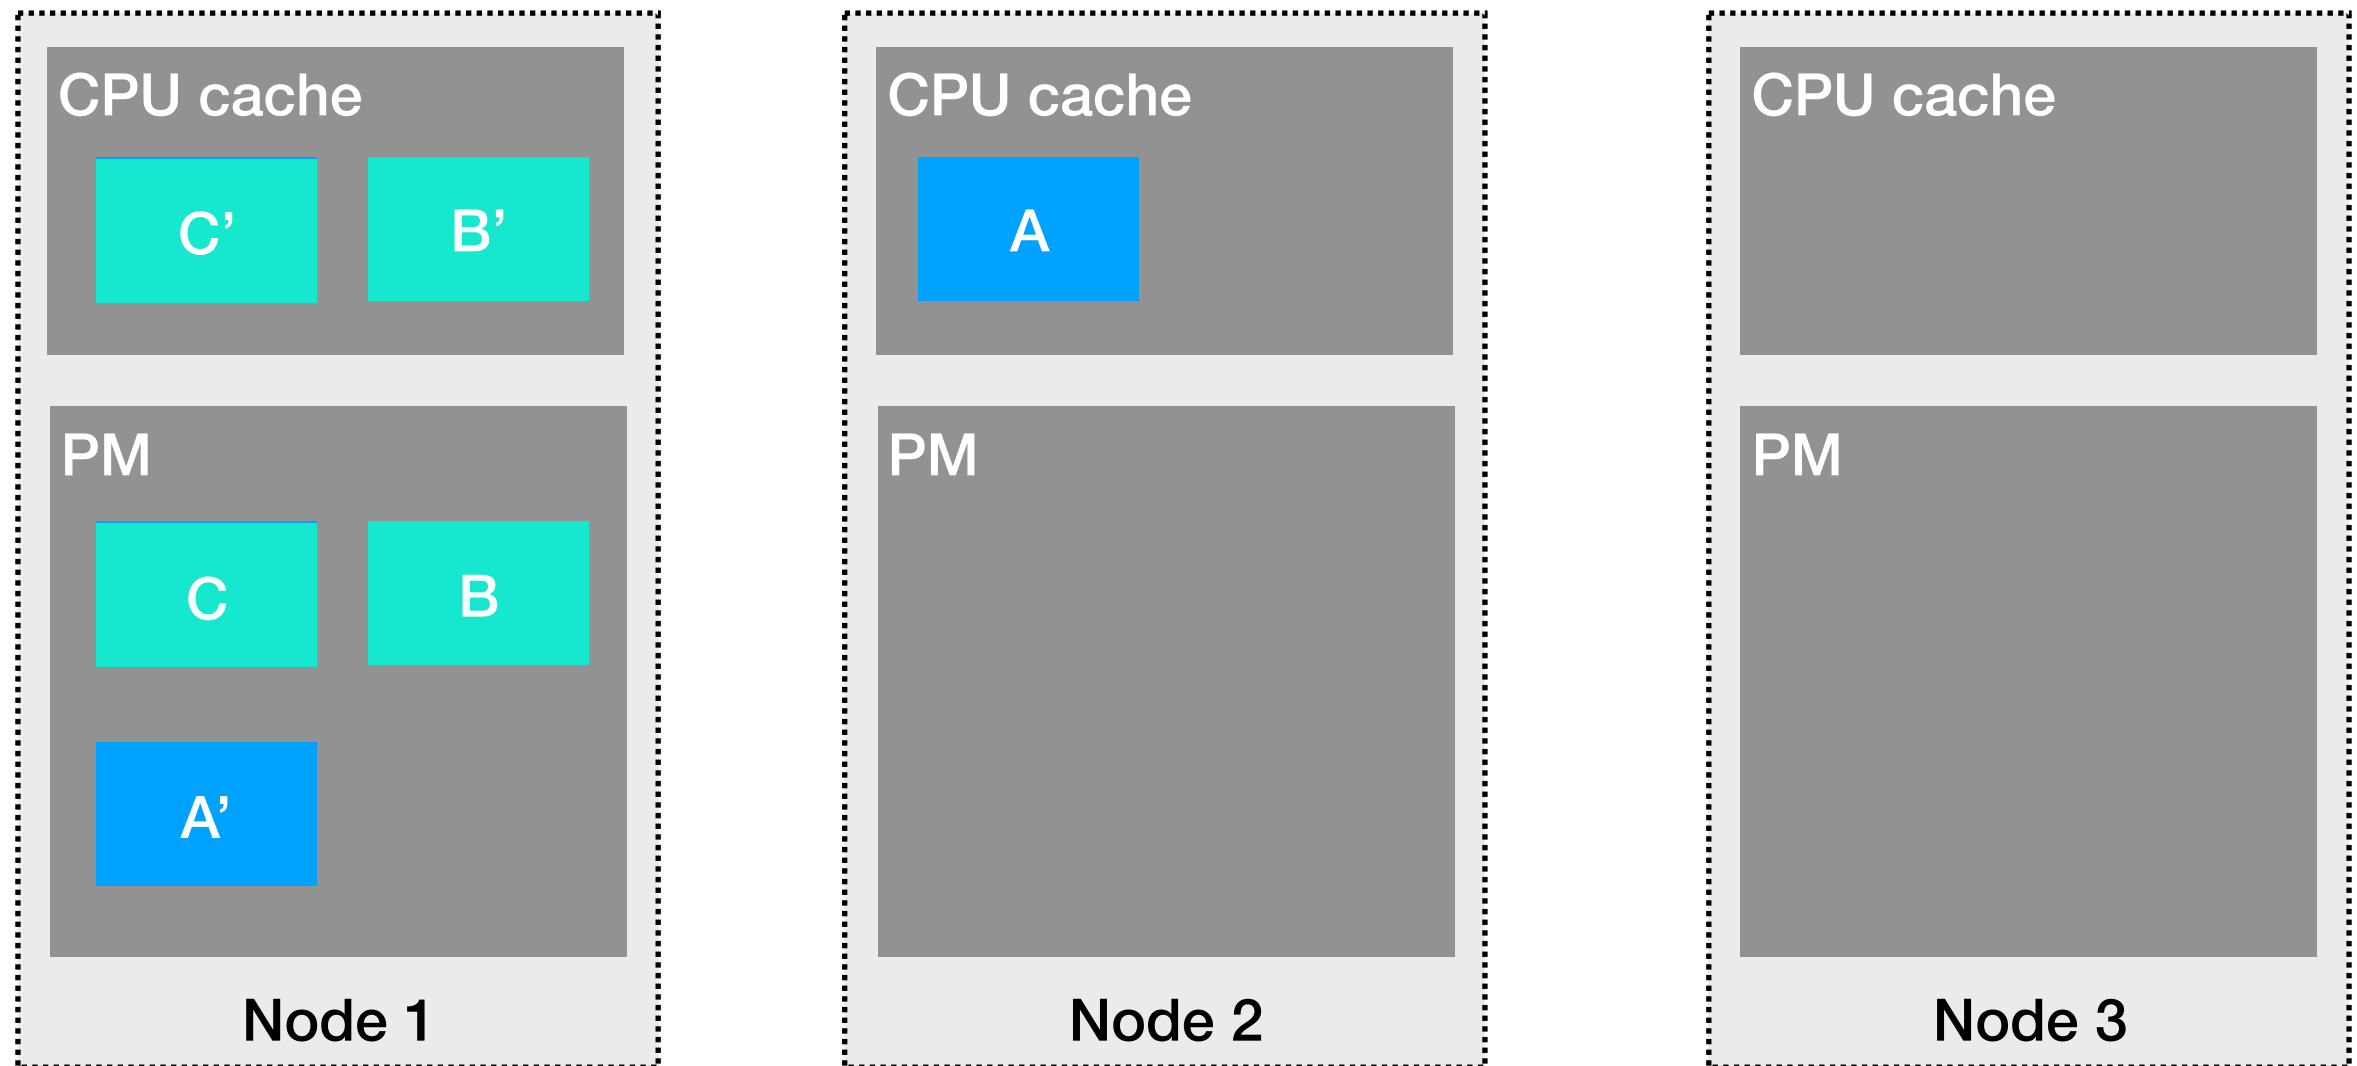
- When to make cached copies coherent?
- When to make data durable and reliability?
- Observations
  - Data-store applications have well-defined **commit points**
  - Commit points: time to make data persistent
  - Visible to storage devices => visible to other nodes



**Exploit application behavior:**

**Make data coherent only at commit points**

# Commit Point



- durable
- coherent
- reliable
- single-node and distributed consistency
- two consistency modes: single/multiple writer


# Flexible Coherence Levels

- Multiple Reader Multiple Writer (***MRMW***)
  - Allows multiple concurrent dirty copies
  - Great parallelism, but weaker consistency
  - ***Three-phase*** commit protocol
- Multiple Reader Single Writer (***MRSW***)
  - Allows only one dirty copy
  - Trades parallelism for stronger consistency
  - ***Single phase*** commit protocol

# MongoDB Results

- Modify MongoDB with ~120 LOC, use MRMW mode
- Compare with *tmpfs*, *PMFS*, *Mojim*, *Octopus* using *YCSB*

Workload	Read	Update	Scan	Insert	R U
A	50%	50%	-	-	-
B	95%	5%	-	-	-
C	100%	-	-	-	-
D	95%	-	-	5%	-
E	-	-	95%	5%	-
F	50%	-	-	-	50%



# Conclusion

- One layer approach: challenges and benefits
- Hotpot: a kernel-level RDMA-based DSPM system
- Hide complexity behind simple abstraction
- Calls for attention to use PM in datacenter
- **Many open problems in distributed PM!**



# Thank You Questions?

Get Hotpot at: <https://github.com/WukLab/Hotpot>

@WukLab

[wuklab.io](http://wuklab.io)

