

Disaggregating Memory with Software-Managed Virtual Cache

Yizhou Shan, Yiyang Zhang



@WukLab



Resource Disaggregation:

**Breaking monolithic
servers into network-
attached, independent
hardware components**


Traditional OSes



- Manages single node and all hardware resources in it
- Bad for hardware heterogeneity and hotplug
- Does not handle component failure

**When hardware is
disaggregated,
the OS
should be also!**

LegO: the *First* Disaggregated OS




OS


**Process
Mgmt**

**Virtual
Memory
System**

**File &
Storage
System**

Network





Key Challenge: Cost of Crossing Network

	Bandwidth	Latency
Mem Bus	50-100 GB/s	~50ns
PCIe 3.0 (x16)	16 GB/s	~700ns
InfiniBand (EDR)	12.5 GB/s	500ns
InfiniBand (HDR)	25 GB/s	<500ns
GenZ	32-400 GB/s	<100ns


- Network **hardware** is much faster than before
- Current network still slower than local memory bus

Observations in Memory Access and Hardware Trends

- Total memory footprint can be large,
- but most accesses go to a small portion
 - 90% => 7MB / 9.6GB (pagerank)
 - 95% => 300MB / 9.6GB
- Faster, smaller memory close to CPU
 - HBM, 3D-stacked
- More computation power at memory
 - PIM/PNM

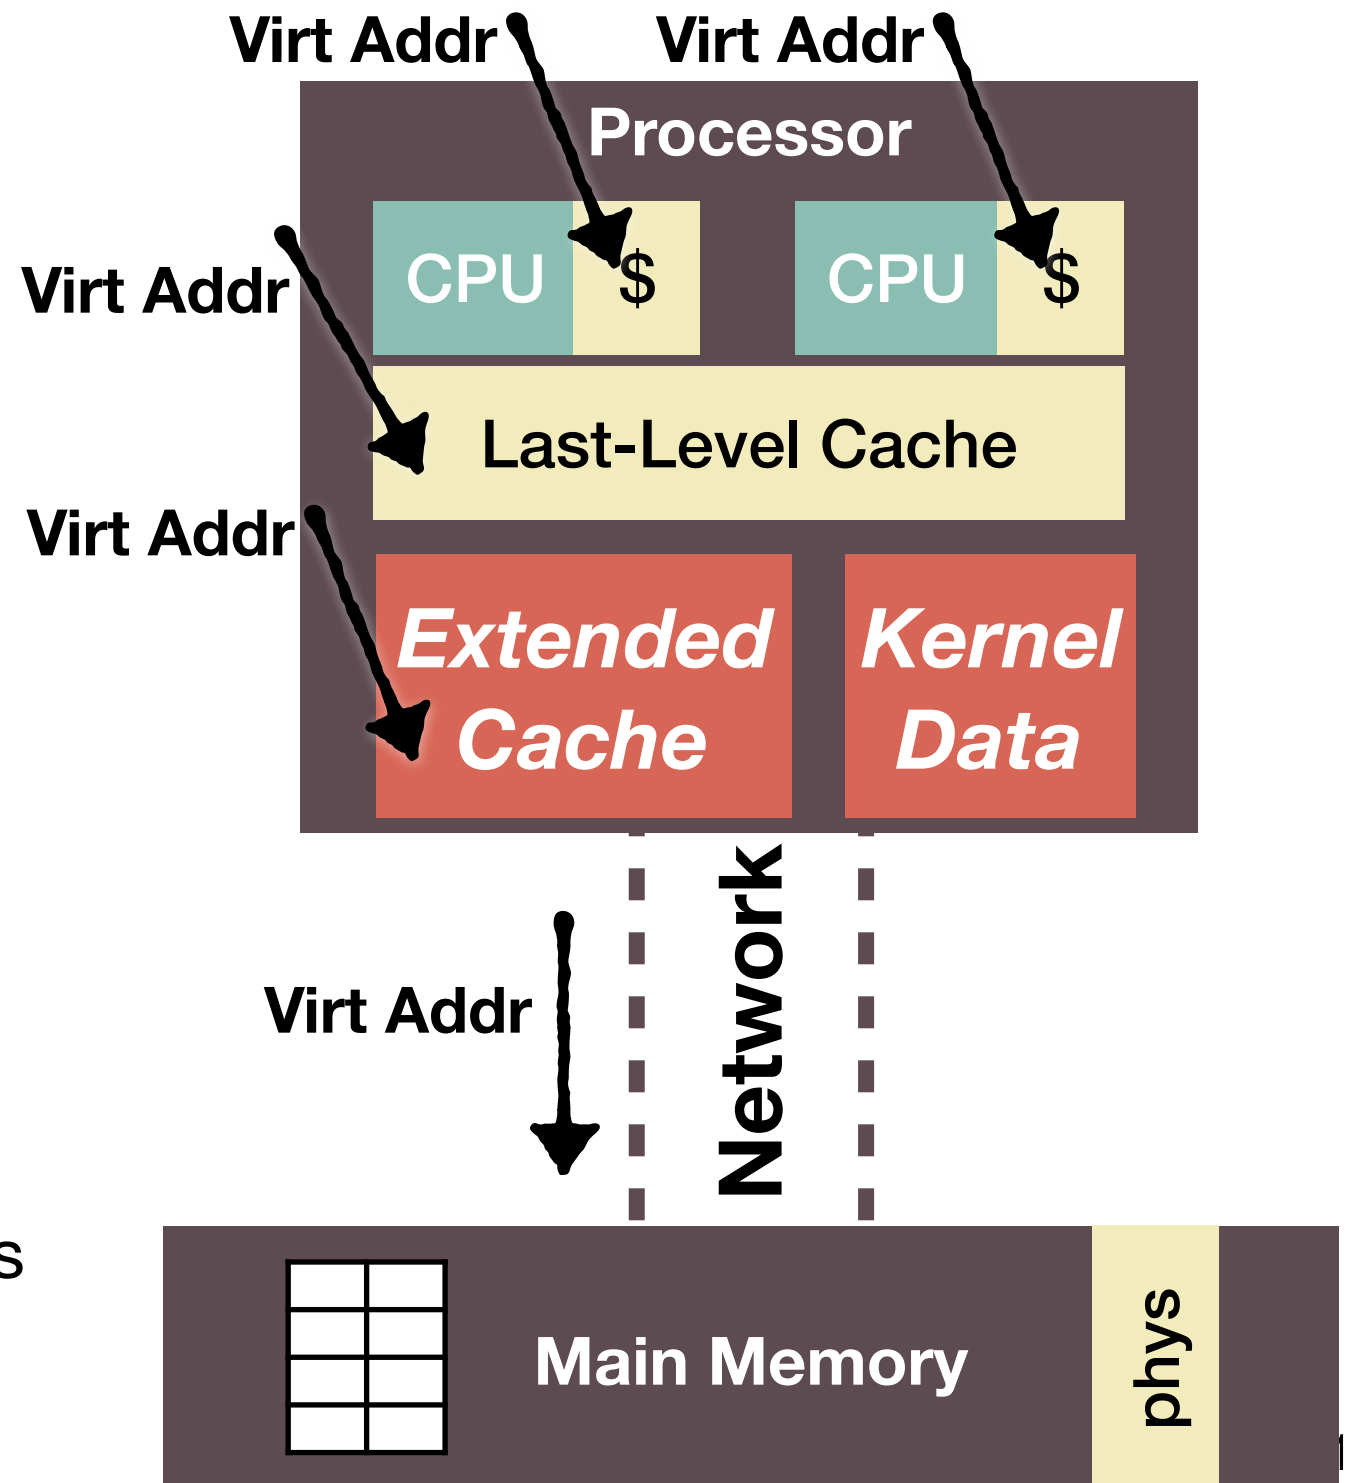
Our Solution: Separate Memory Perf and Capacity

- **Bigger memory behind network**
- **Extended cache at processor**
 - HBM or regular DRAM
 - Can be software-managed
- Separate physical mem for kernel



Clean Separation of Processor and Memory Functionalities

- Important for heterogeneity, flexibility, and failure independence
- Memory components manage
 - Virtual and physical memory spaces
 - Virtual to physical memory mapping
- Processors
 - Only see virtual memory addresses
 - Software-managed virtual cache



Other Challenges


- Handling component failure
- Manage distributed, heterogeneous resources
- Fitting micro-OS services in hardware controller
- Implementing Lego on current servers

Implementation and Emulation

- Lego built from scratch, >200K LOC and growing
 - Runs all Linux ABIs and unmodified binaries
 - Manages disaggregated processor, memory, storage
 - Global resource manager
- Hardware emulation
 - Use regular servers to emulate hardware devices
 - DRAM organized as extended cache, managed by page fault handler

Preliminary Results

- One processor, one memory, one storage, connected with 40Gbps InfiniBand
- **Phoenix** [1]: single-node MapReduce implementation, running word count of 2GB file



Conclusion

- Resource disaggregation calls for new system
- **Lego**: new OS designed and built for datacenter resource disaggregation
- Separating memory performance and capacity, and processor and memory functionalities
- Many challenges and many potentials

Thank you Questions?

@WukLab

wuklab.io



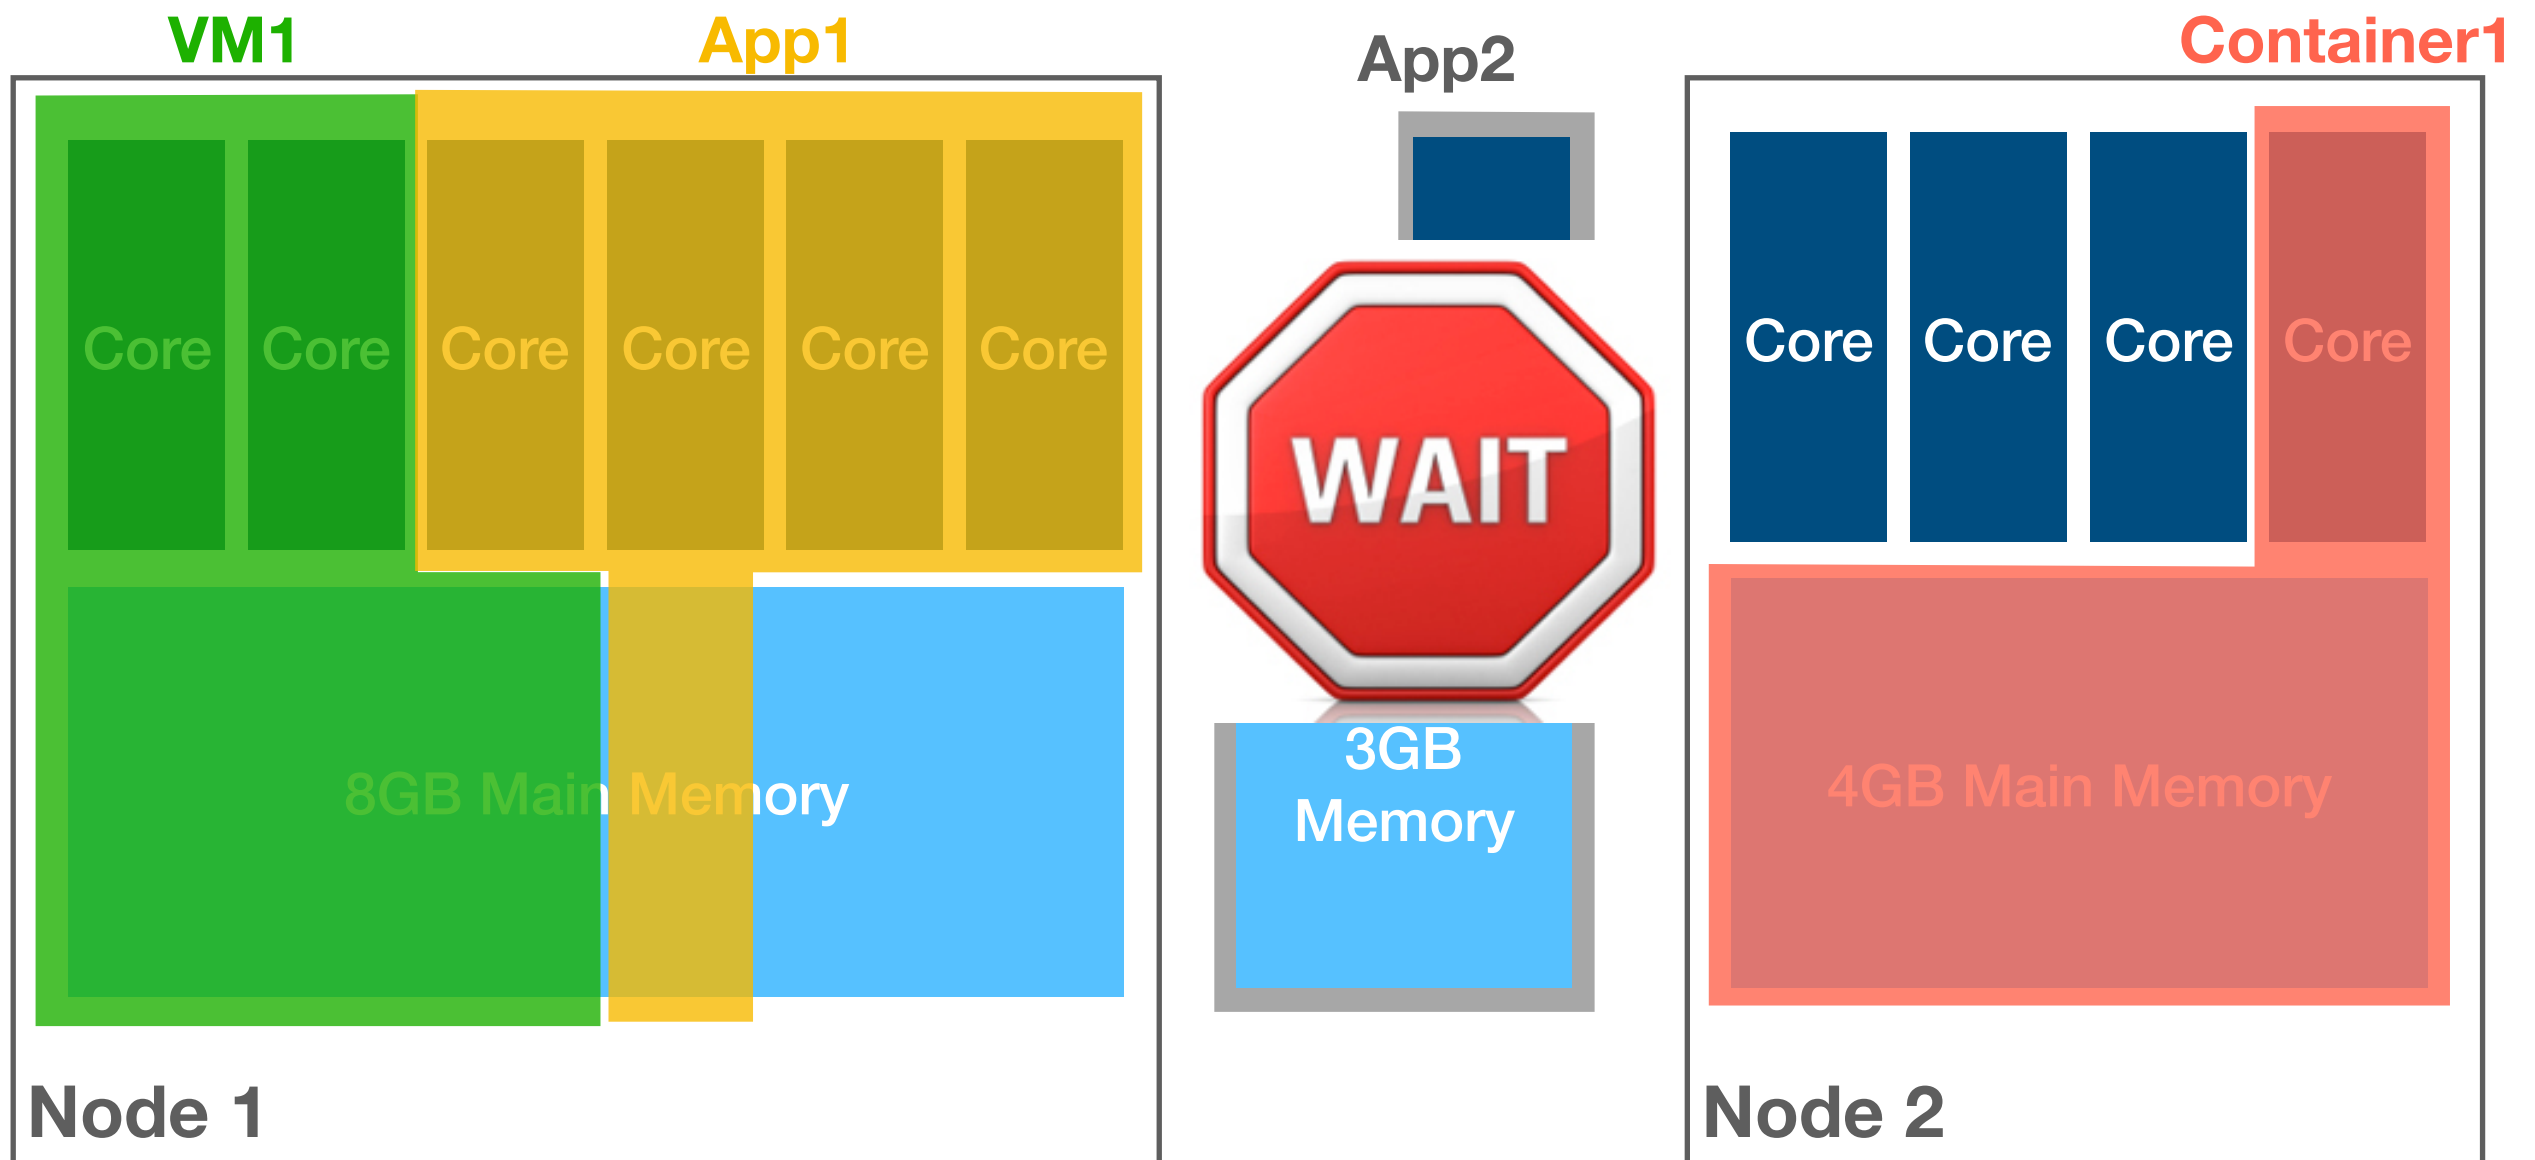
Split Container: Running Containers beyond Physical Machine Boundaries

Yilun Chen, Yiying Zhang

@WukLab




Resource Allocation in Datacenters




Physical Machine Boundary!

CPU/Memory Usages across Machines and across Jobs



99th-to-Median Memory Usage Across Machines


Source: Gu et al. "Efficient Memory Disaggregation with Infiniswap" NSDI'17



Google Cluster Trace "<https://github.com/google/cluster-data>"

Modern Datacenter Applications Have Heterogeneous CPU/Memory Requirements

Resource Utilization in Production Clusters



* Google Production Cluster Trace Data. "<https://github.com/google/cluster-data>"

Unused Resource + Waiting/Killed Jobs Because of Physical-Node Constraints




Physical Resource Disaggregation

- Great support of heterogeneity
- Very flexible in resource management
- But needs hardware, network, and OS changes

**Is there any less disruptive
way to achieve better
resource utilization and
elasticity?**

Virtually Disaggregated Datacenter

- Use resources on remote (distributed) machines



Using Remote/Distributed Resources

- Was a popular idea in 90s
 - Remote memory/paging/swap
 - Network block device
 - Distributed shared memory (DSM)
- No production-scale adoption
 - Cost of network communication
 - Coherence traffic



REVISIT YOUR
OLD IDEAS
WITH NEW EYES.

Remote/Distributed Memory in Modern Times

- New application trends
 - Large parallelism
 - New computation and memory requirements
 - New programming models
- Network is 10x-100x faster
 - InfiniBand: 200Gbps, <500ns
 - GenZ: 32-400GB/s, <100ns

Recent New Attempts

- Distributed Shared Memory
 - Grappa
 - Hotpot (Distributed Shared Persistent Memory)
- Network swapping
 - InfiniSwap
- Non-coherent distributed memory
 - VMware
- How to communicate across nodes?
At what level of transparency?

Message Passing

- New programming languages
 - Use message passing instead of shared memory to do thread communication
 - Golang channel, Erlang actors, Akka library
- New application development practices
 - Use multiple processes instead of multithreading
 - Use message passing instead of shared memory
 - Nginx, Apache Web Server, Node.js

Splitting Containers with Message Passing

- Splitting a container across physical machines
 - Both processes and memory
 - Elastic, good resource utilization
- Use message passing for both inter- and intra-process communication
 - No coherence traffic
 - Easy to track and optimize inter-node communication

Challenges in Splitting Containers

- Performance overhead of crossing network
- Performance imbalance
- Management of split resources
- QoS
- Failure handling

Conclusion

- Splitting computation and memory not a new idea
- But new application and network properties
- Splitting container across physical machines
- Use message passing for all inter- and intra-process communication
- More challenges and opportunities

Thank you Questions?

@WukLab

wuklab.io



Mitsume: an Object-Based Remote Memory System

Shin-Yeh Tsai, Yiying Zhang

@WukLab



One-Sided Remote Memory/NVM

- One-sided devices
 - Devices without (general) computation power
 - Can only be read and written to with limited, low-level APIs
 - Disaggregated memory
 - NVMe over fabrics
- Cheap, low energy


Remote Memory Challenges and Opportunities

- Challenges
 - No computation power at memory
 - Remote memory can fail independently
- Opportunities
 - Can trace and control (remote) memory accesses
 - No local accesses that can violate atomicity


**How to use one-sided remote
memory/NVM devices?**

Local Multiprocessor Shared Memory

Single Node



Remote Memory



Remote Memory and Local Memory Comparison

- Similarities
 - No computation power
 - Multiple processors (cores) can read and write to
- Differences (remote memory)
 - No hardware coherence
 - Can fail independently (and more often)
 - Larger but slower than local memory

Our View of Remote Memory


- Treat remote memory as a raw data-store hardware
 - Similar to DRAM chips in local main memory
 - Fast, cheap data store
- Extract the control and intelligence apart
 - Similar to memory controller in local main memory
 - But managed in software

Mitsume*: an Object-Based Remote Memory System

- Separate data and control path
- Data: one-sided
 - Client nodes read/write to remote memory
 - Multiple processors (cores) can read and write to
- Control: two-sided
 - Global software controller manages remote memory and talks to clients via two-sided operations

* Mitsume means three eyes in Japanese and is from the manga and anime Mitsume ga Tooru (the Three-Eyed One)

Mitsume Architecture



Mitsume Data Organization

- Data stored and located by “object”
- Updates to an object guaranteed atomic and append-only
- Each object can have multiple versions
- Flexible physical locations of (versions) objects
- Each object can have their own replication factor

Global Control

- Allocate physical memory at remote memory
- Garbage collection
- Ensures QoS for different clients
- Resource management and load balancing
- Failure handling
- Security

Usage Models

- Key-value store
- Version system
- Remote swap
- Messaging system
- Pub/Sub

Conclusion

- One-sided memory/NVM devices are useful
- Learn from local memory system
- Separate data and control of remote memory
- Many usage possibilities of remote memory

Thank you Questions?

@WukLab

wuklab.io

