Kernel:
to be or not to be?

Yiying Zhang

About Me and WukLab

WukLab: Building next-gen datacenter systems

- Fields
  - Operating systems
  - Distributed systems
  - Datacenter networking
  - Computer architecture

"I see myself as a generalist -- I am attracted to the biggest problem I can find, regardless of area"

About Me and WukLab

WukLab: Building next-gen datacenter systems

- Fields
  - Operating systems
  - Distributed systems
  - Datacenter networking
  - Computer architecture

"I see myself as a generalist -- I am attracted to the biggest problem I can find, regardless of area"

But most of our works are in kernel

What should go into kernel?

1983

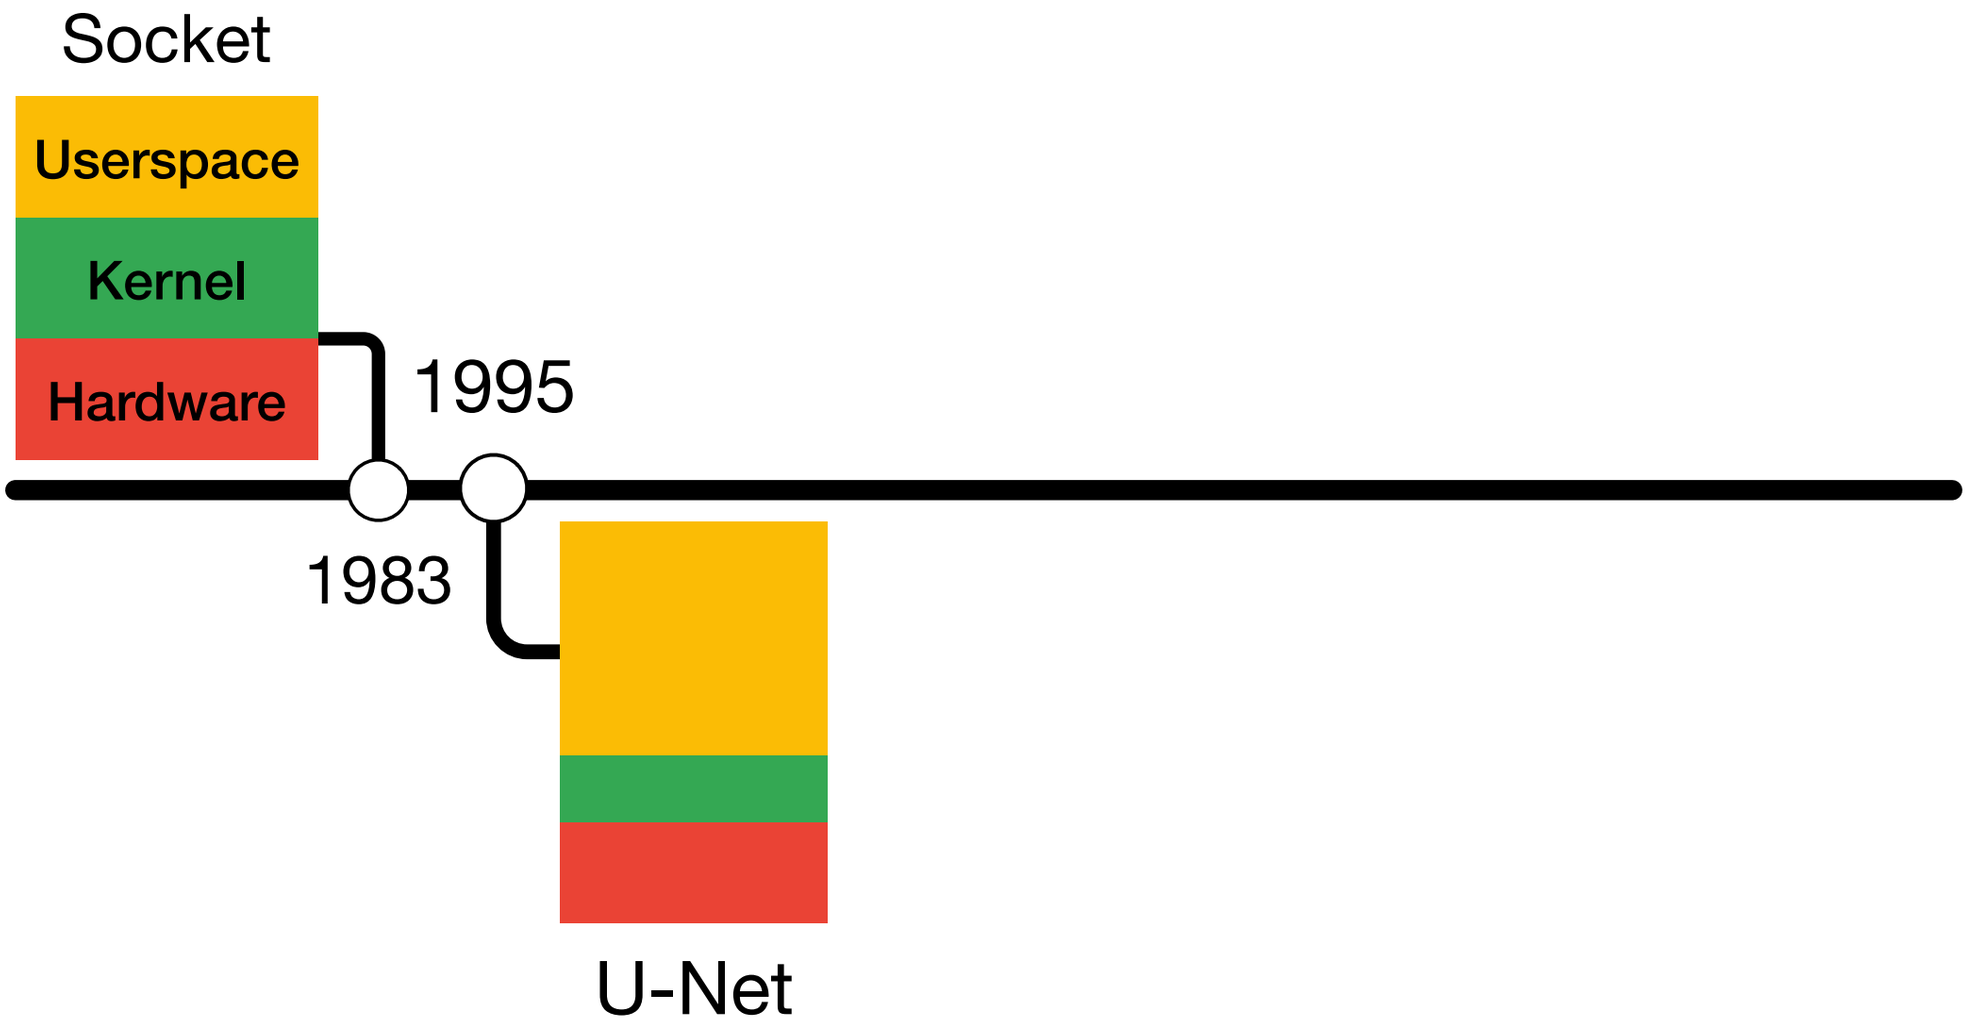
Socket
Userspace
Kernel
Hardware
1995
1983
U-Net

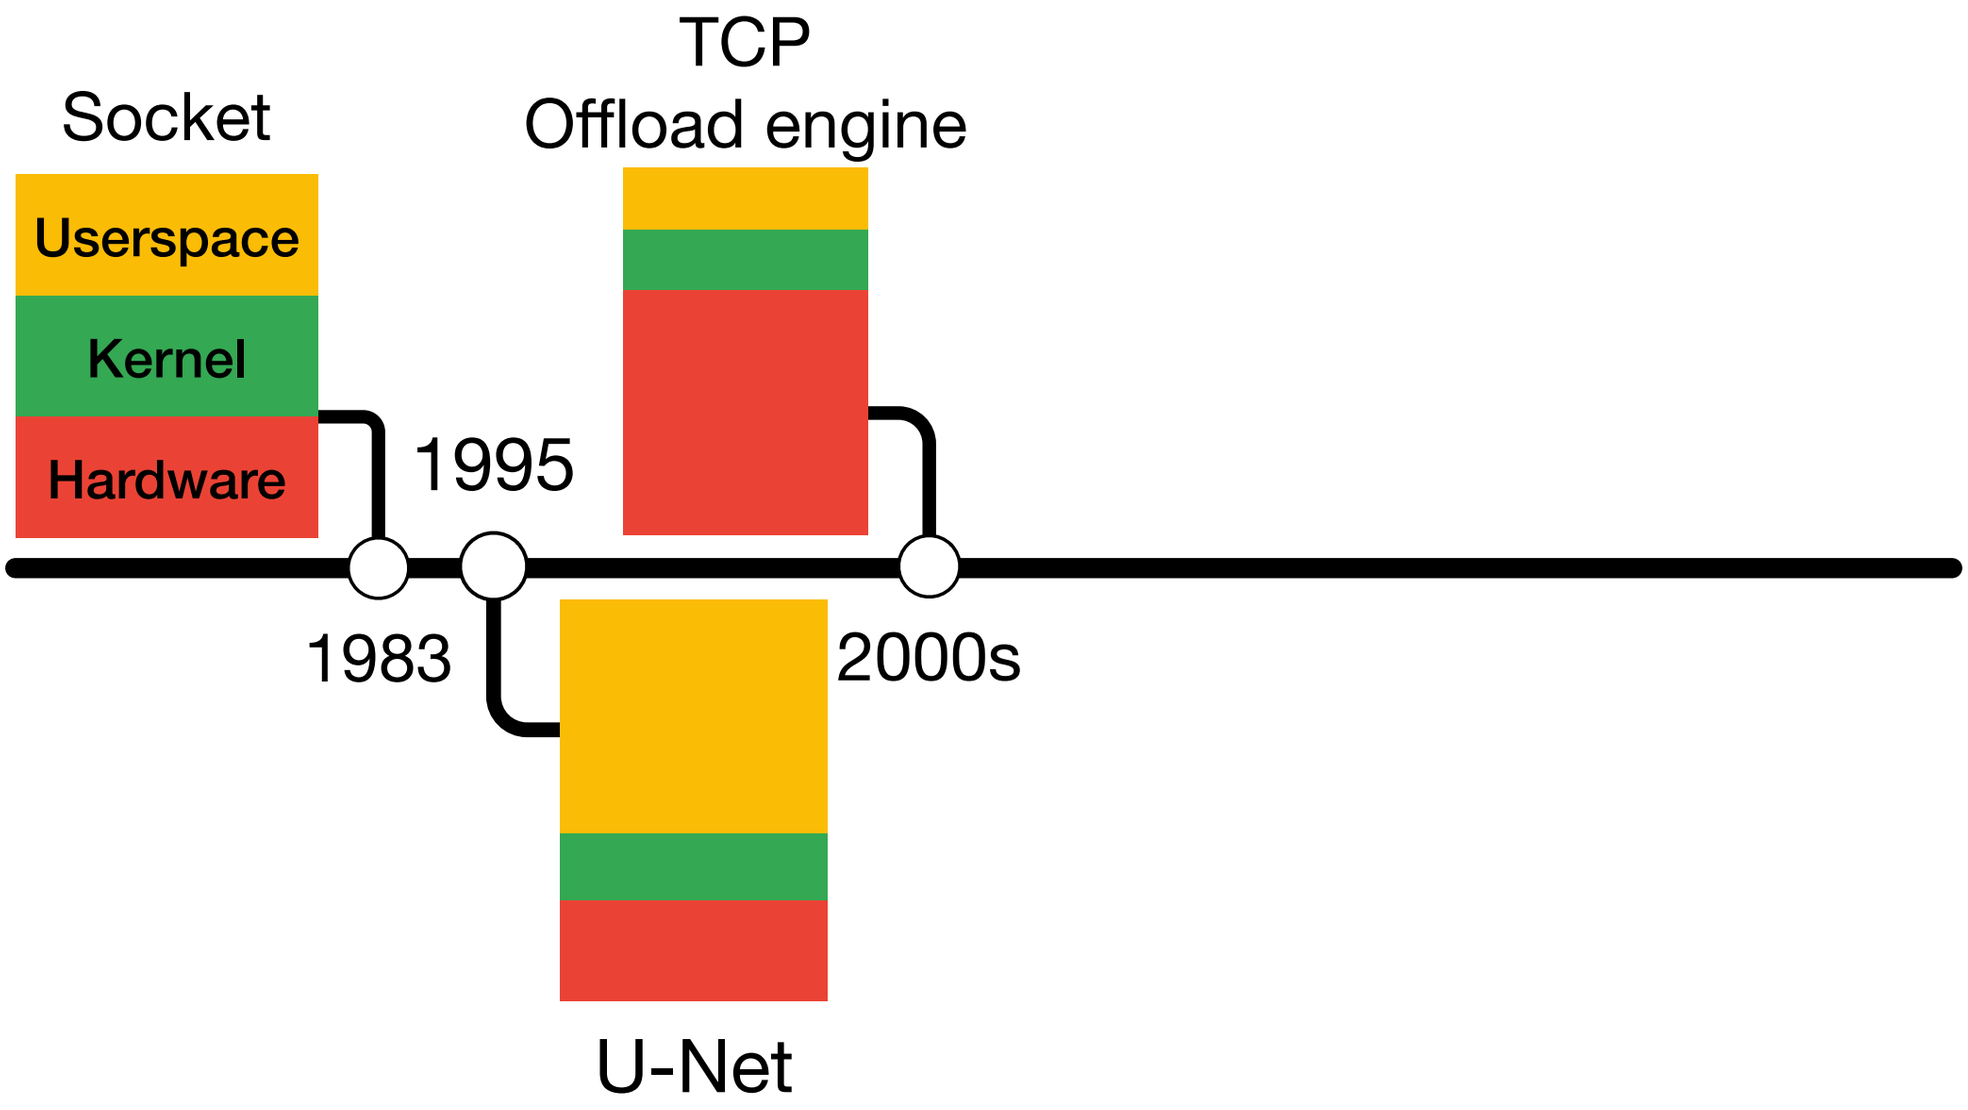
Socket
Userspace
Kernel
Hardware
1983
TCP
Offload engine
1995
2000s
U-Net

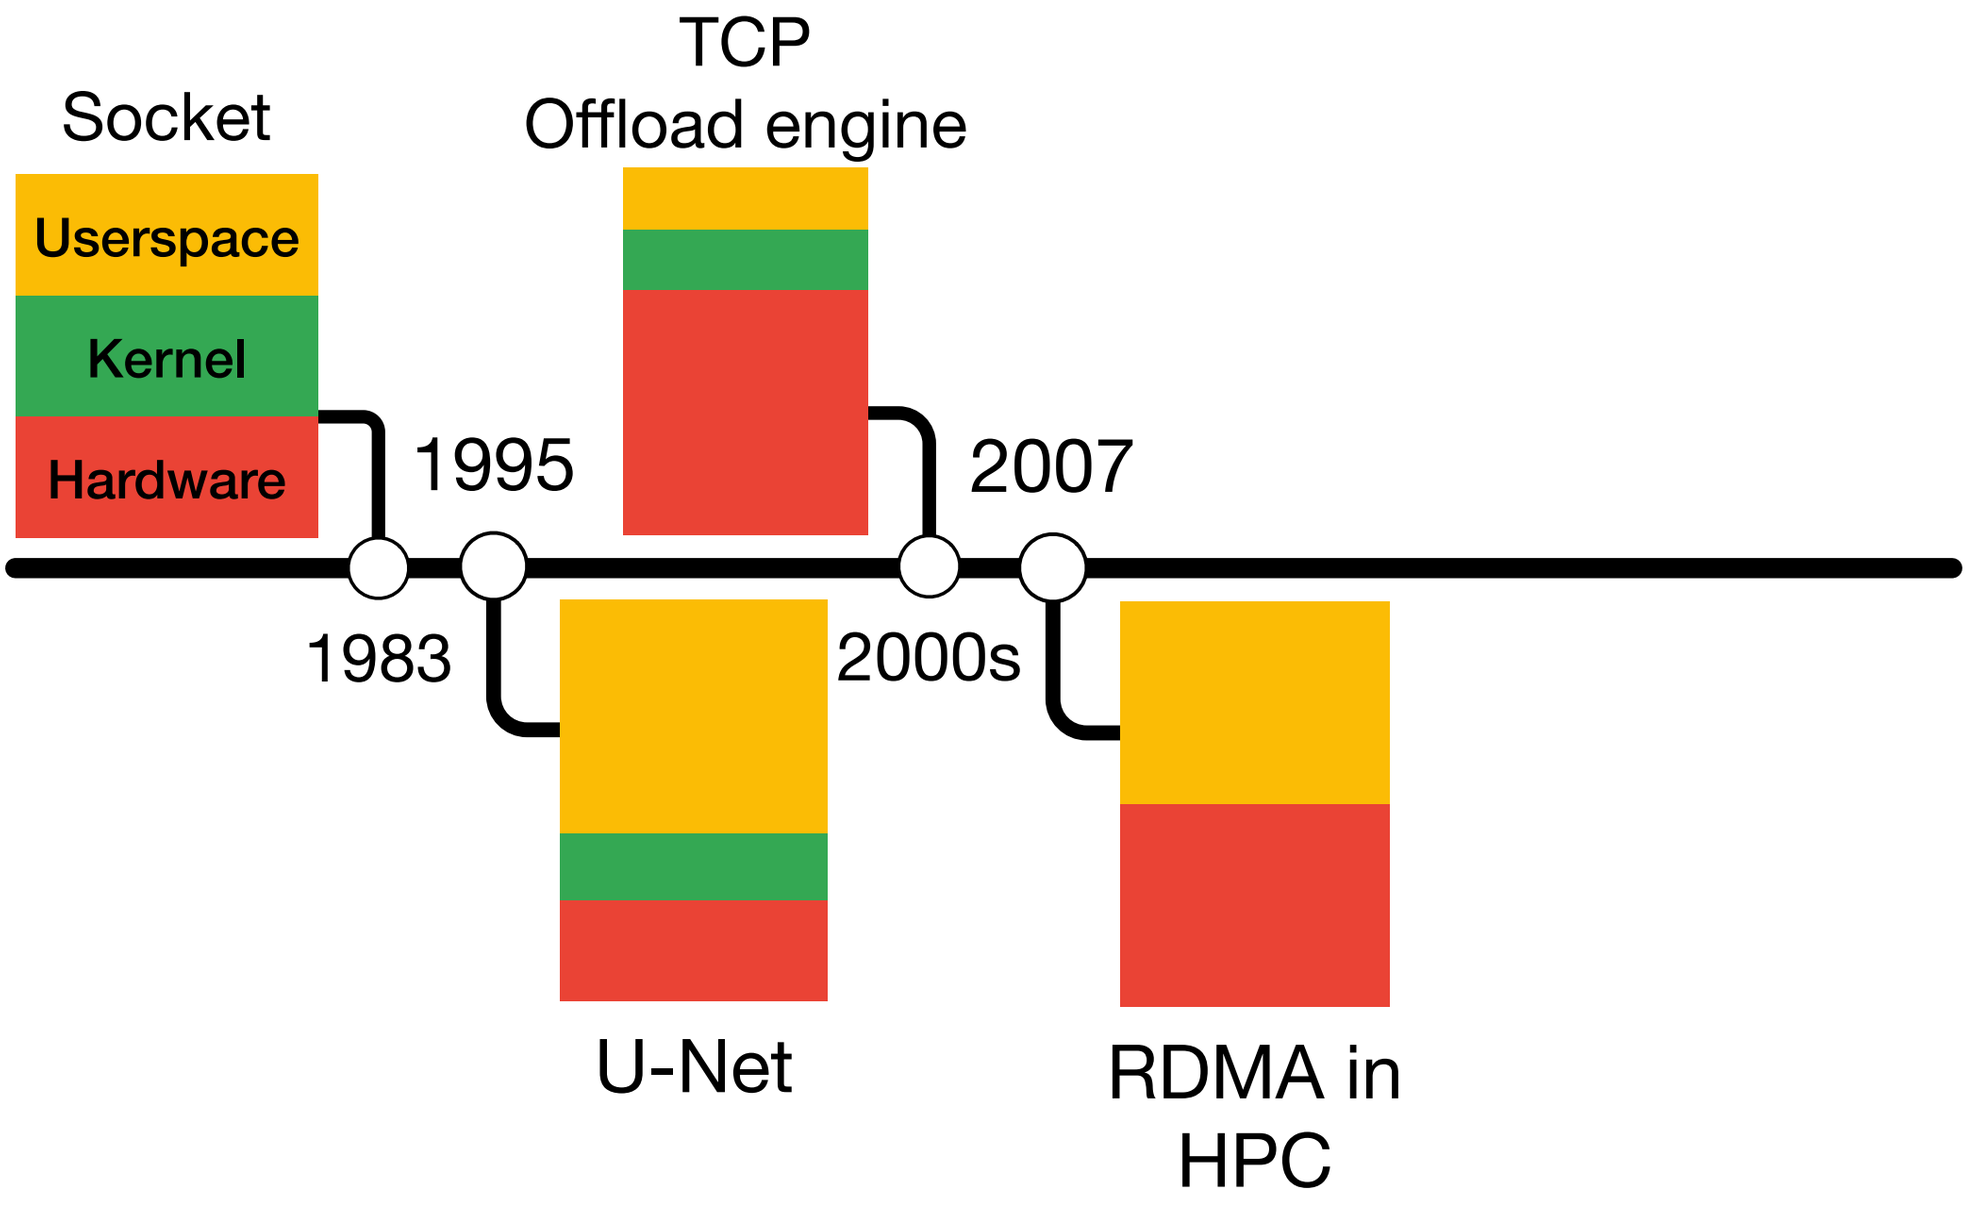
Socket
Userspace
Kernel
Hardware
1995
TCP
Offload engine
2007
1983
U-Net
2000s
RDMA in
HPC

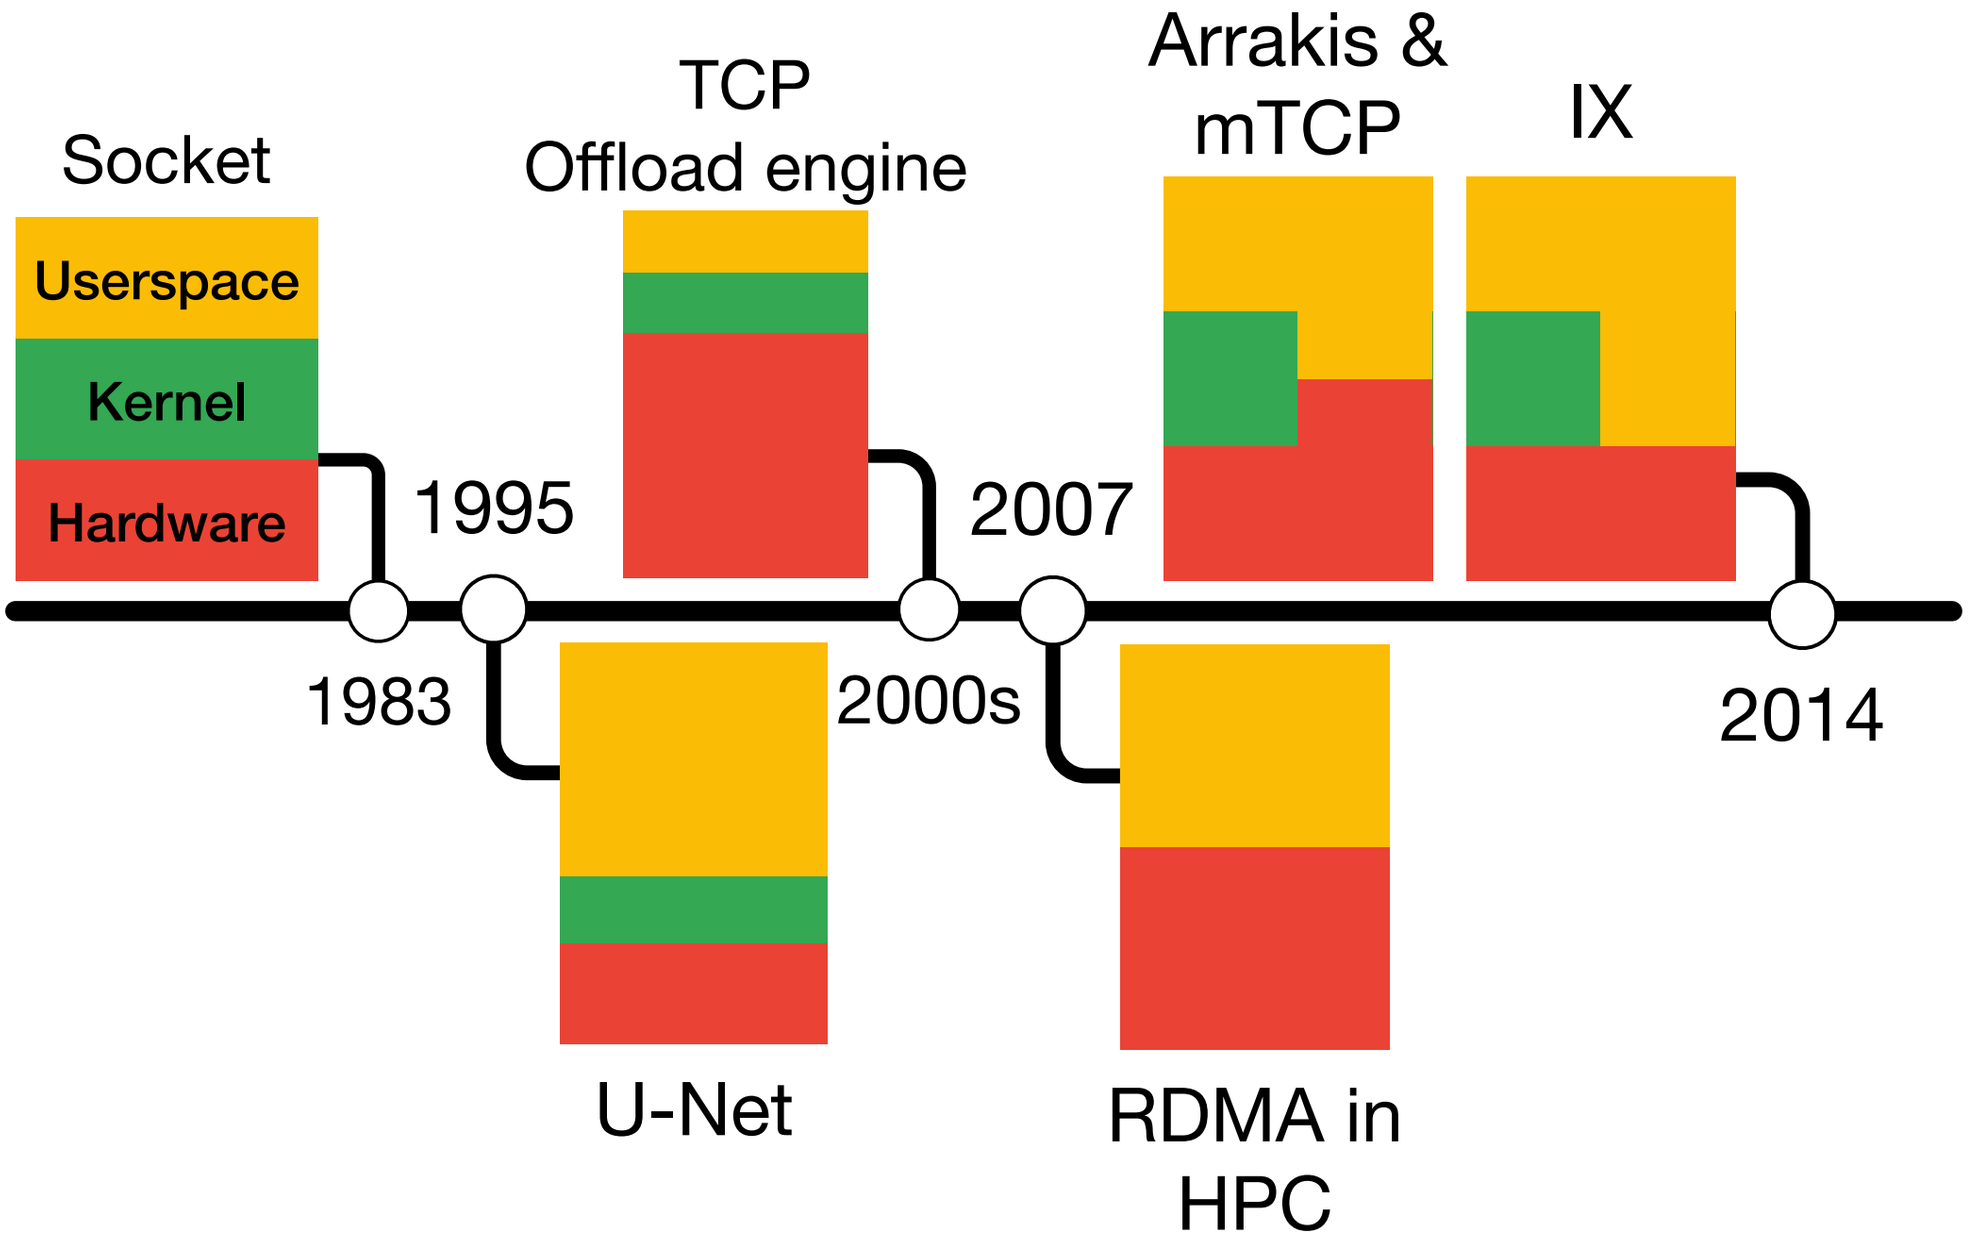
Socket
Userspace
Kernel
Hardware
1995
TCP
Offload engine
2007
Arrakis &
mTCP
IX
1983
2000s
2014
U-Net
RDMA in
HPC

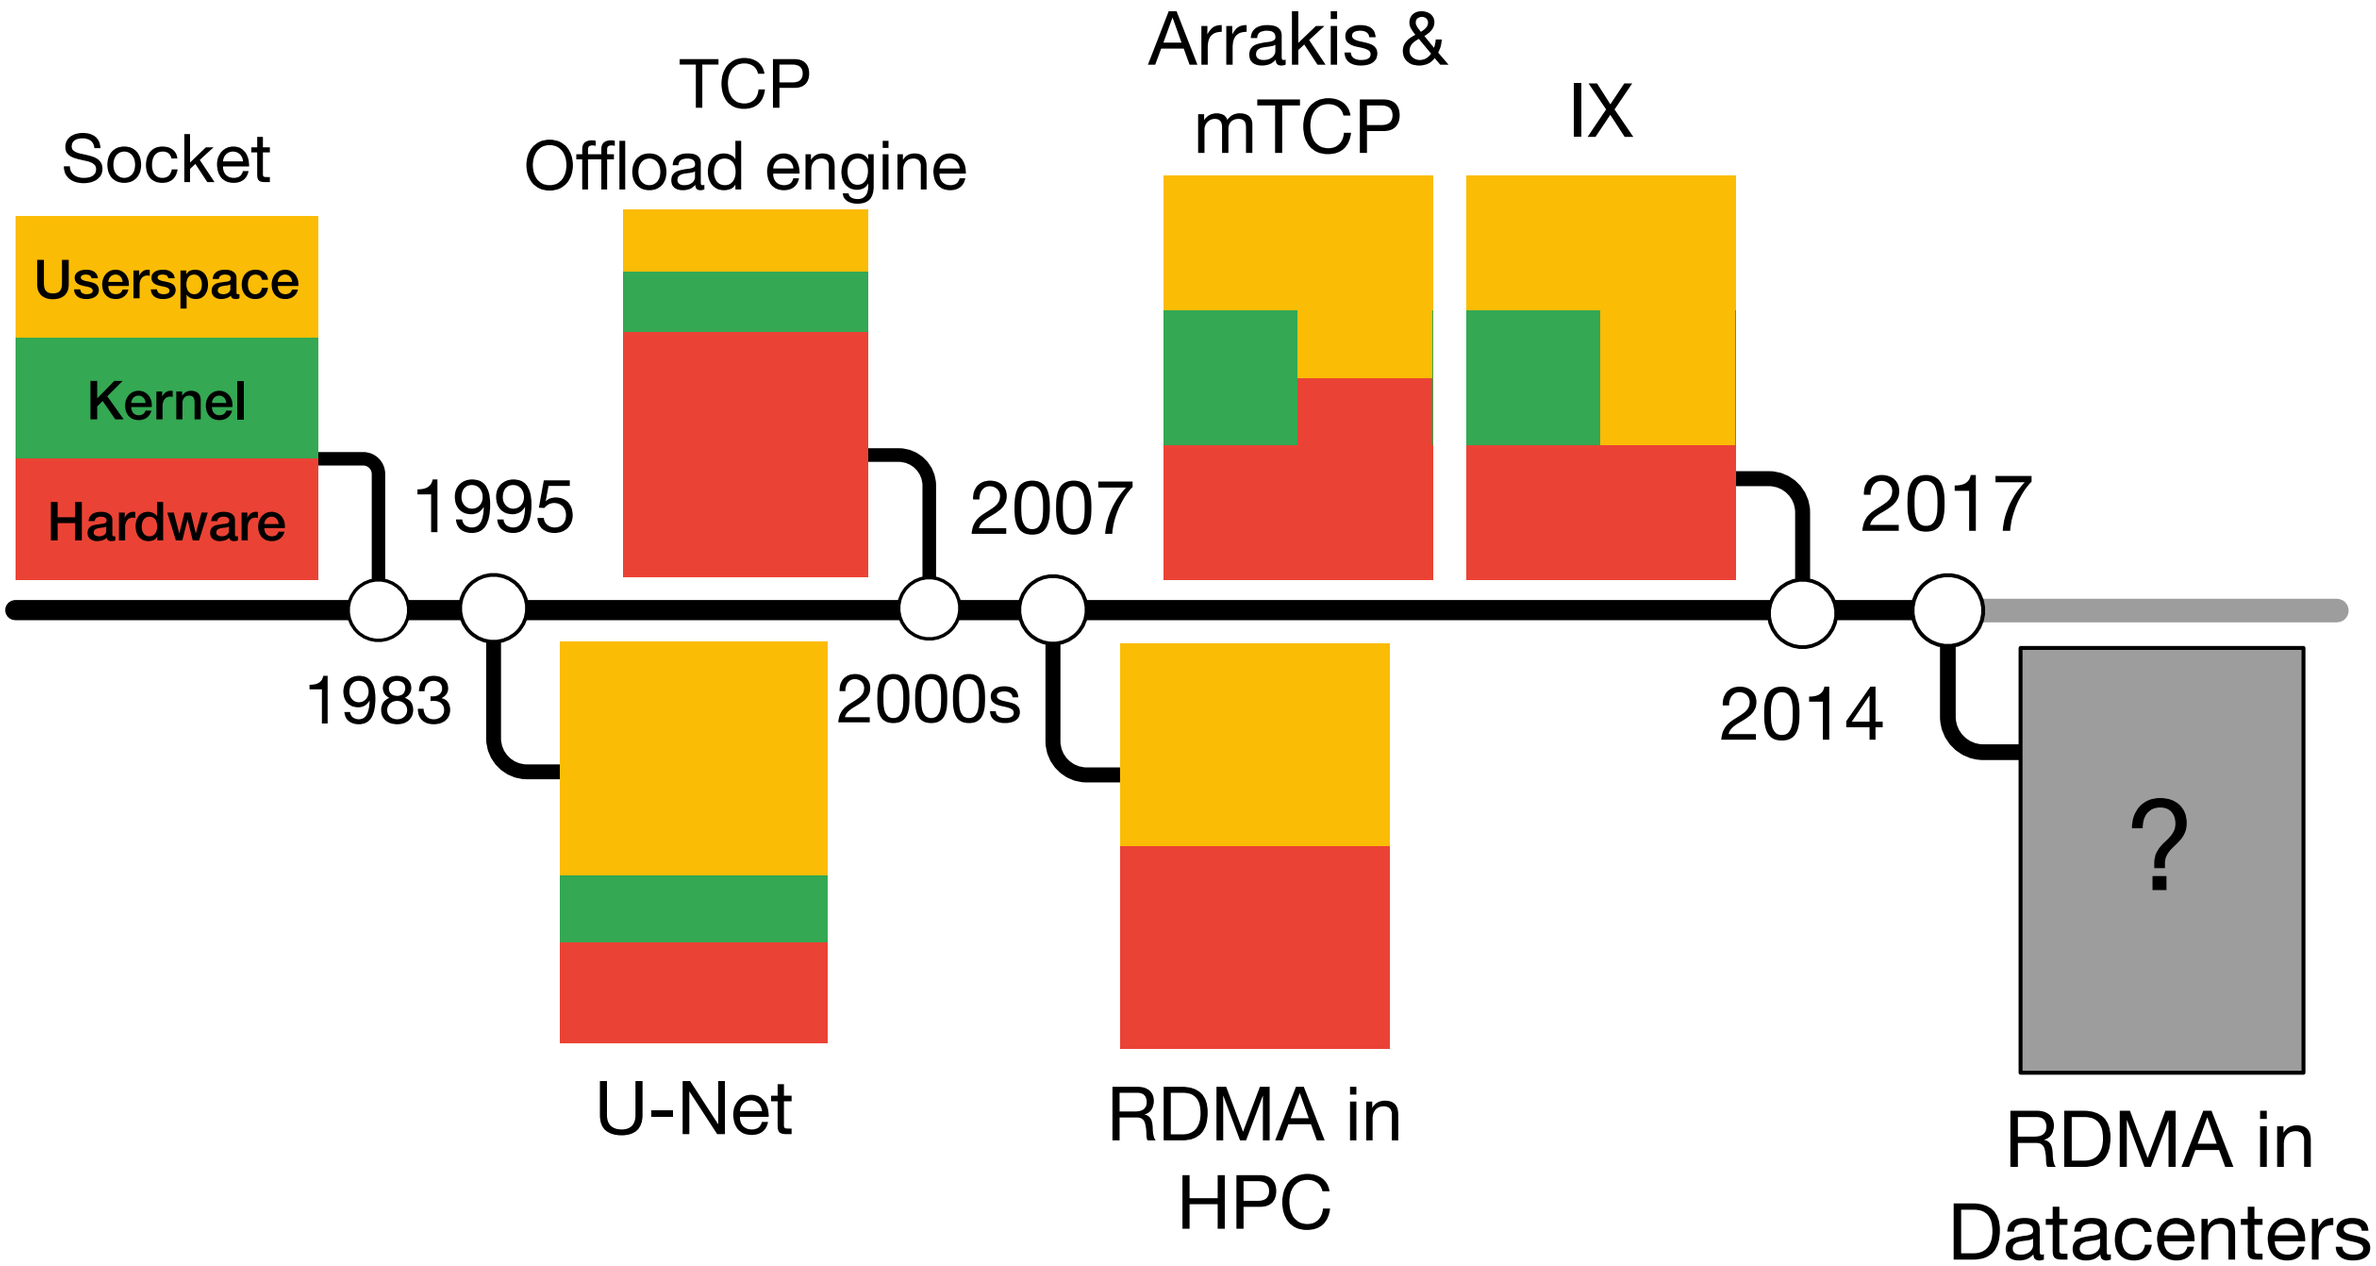
Socket
Userspace
Kernel
Hardware
1995
TCP
Offload engine
2007
Arrakis &
mTCP
IX
2017
1983
2000s
2014
?
U-Net
RDMA in
HPC
RDMA in
Datacenters

Userspace

Hardware

RDMA

RDMA

Userspace

Hardware

User Space

Kernel Space

Hardware

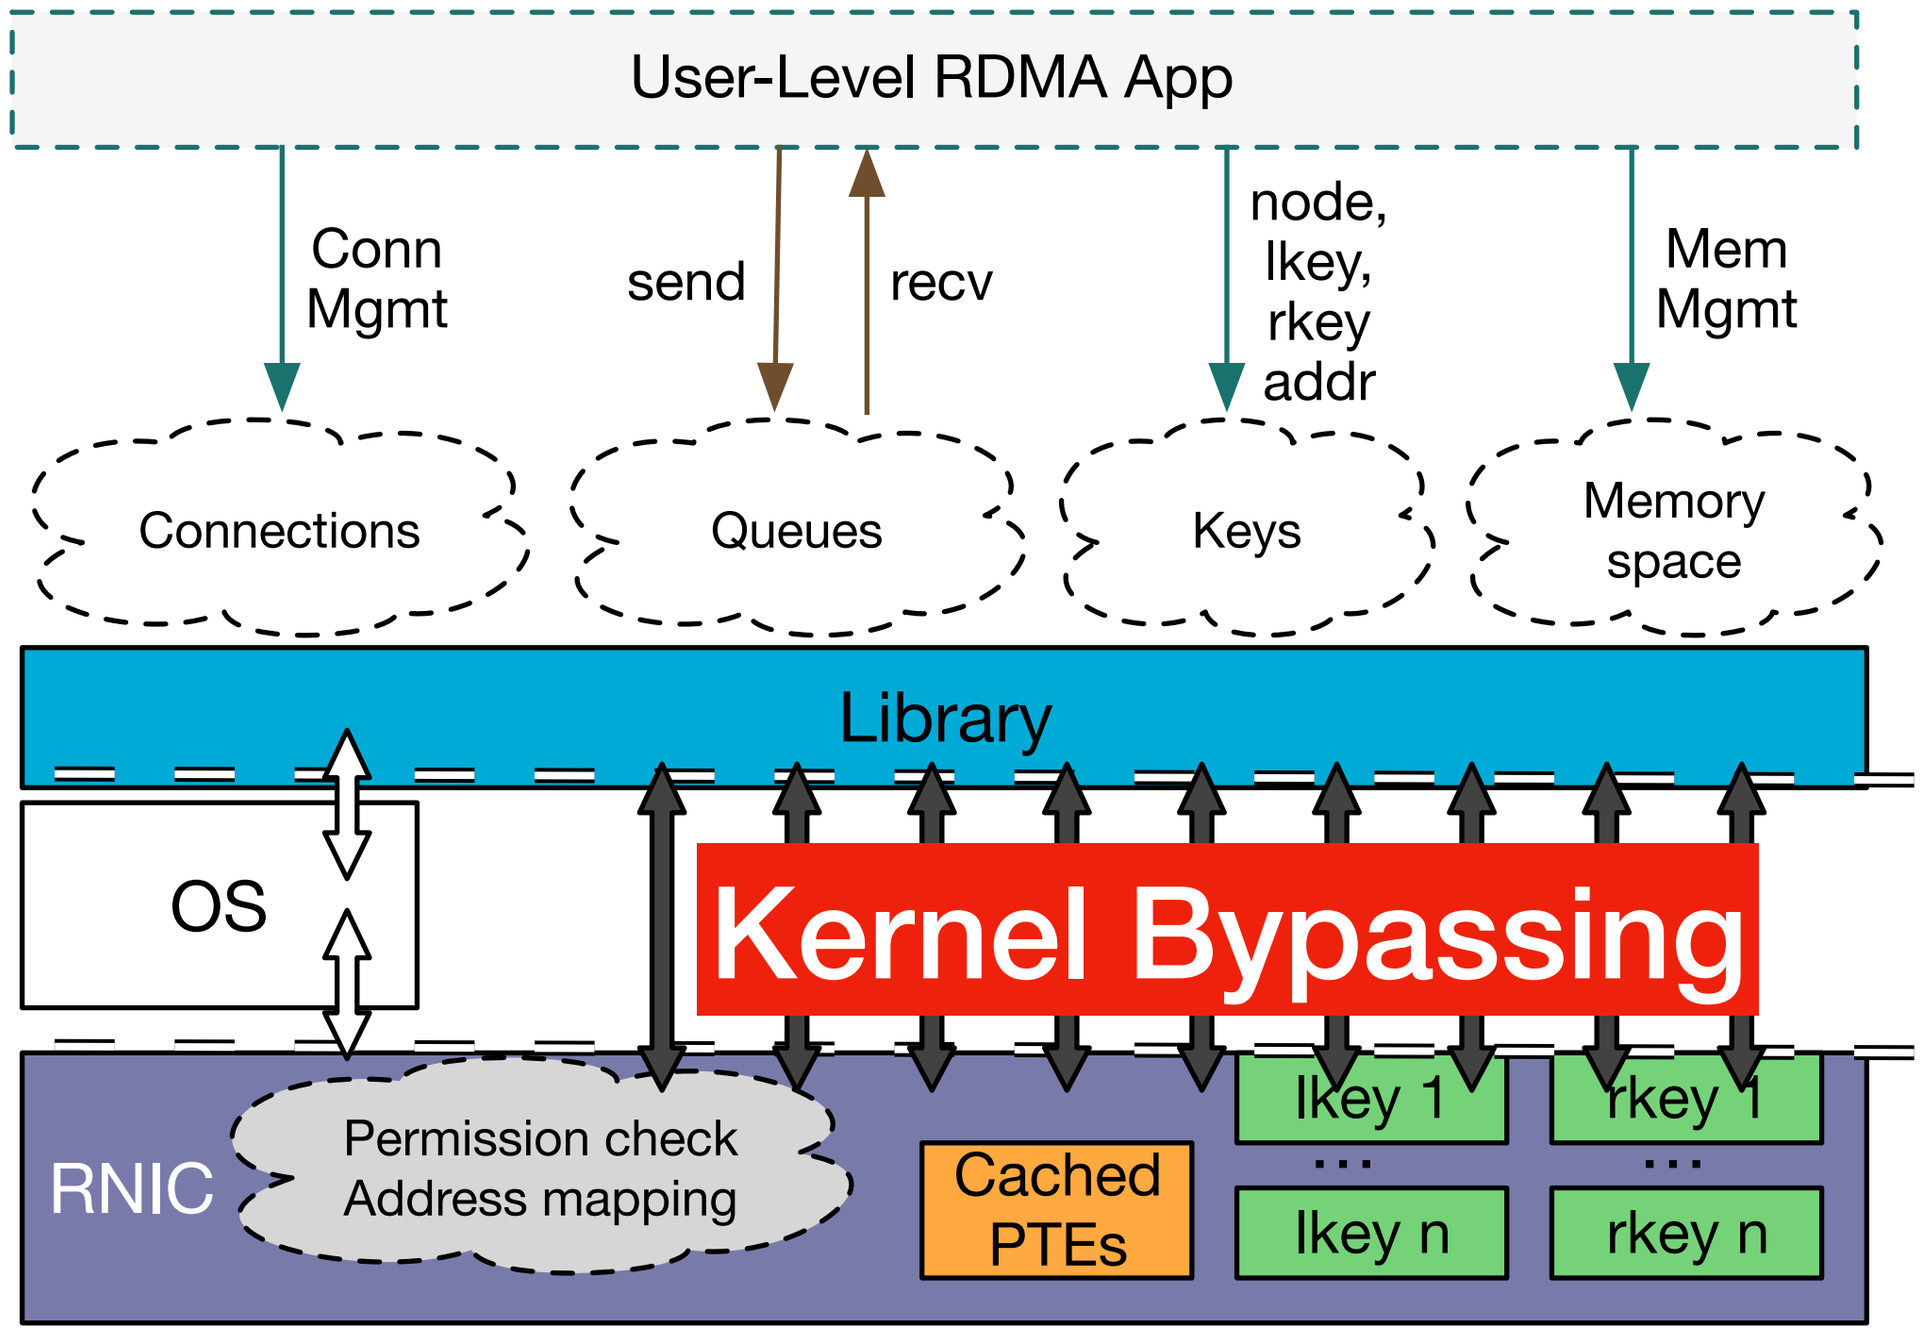

Userspace

Hardware

RDMA

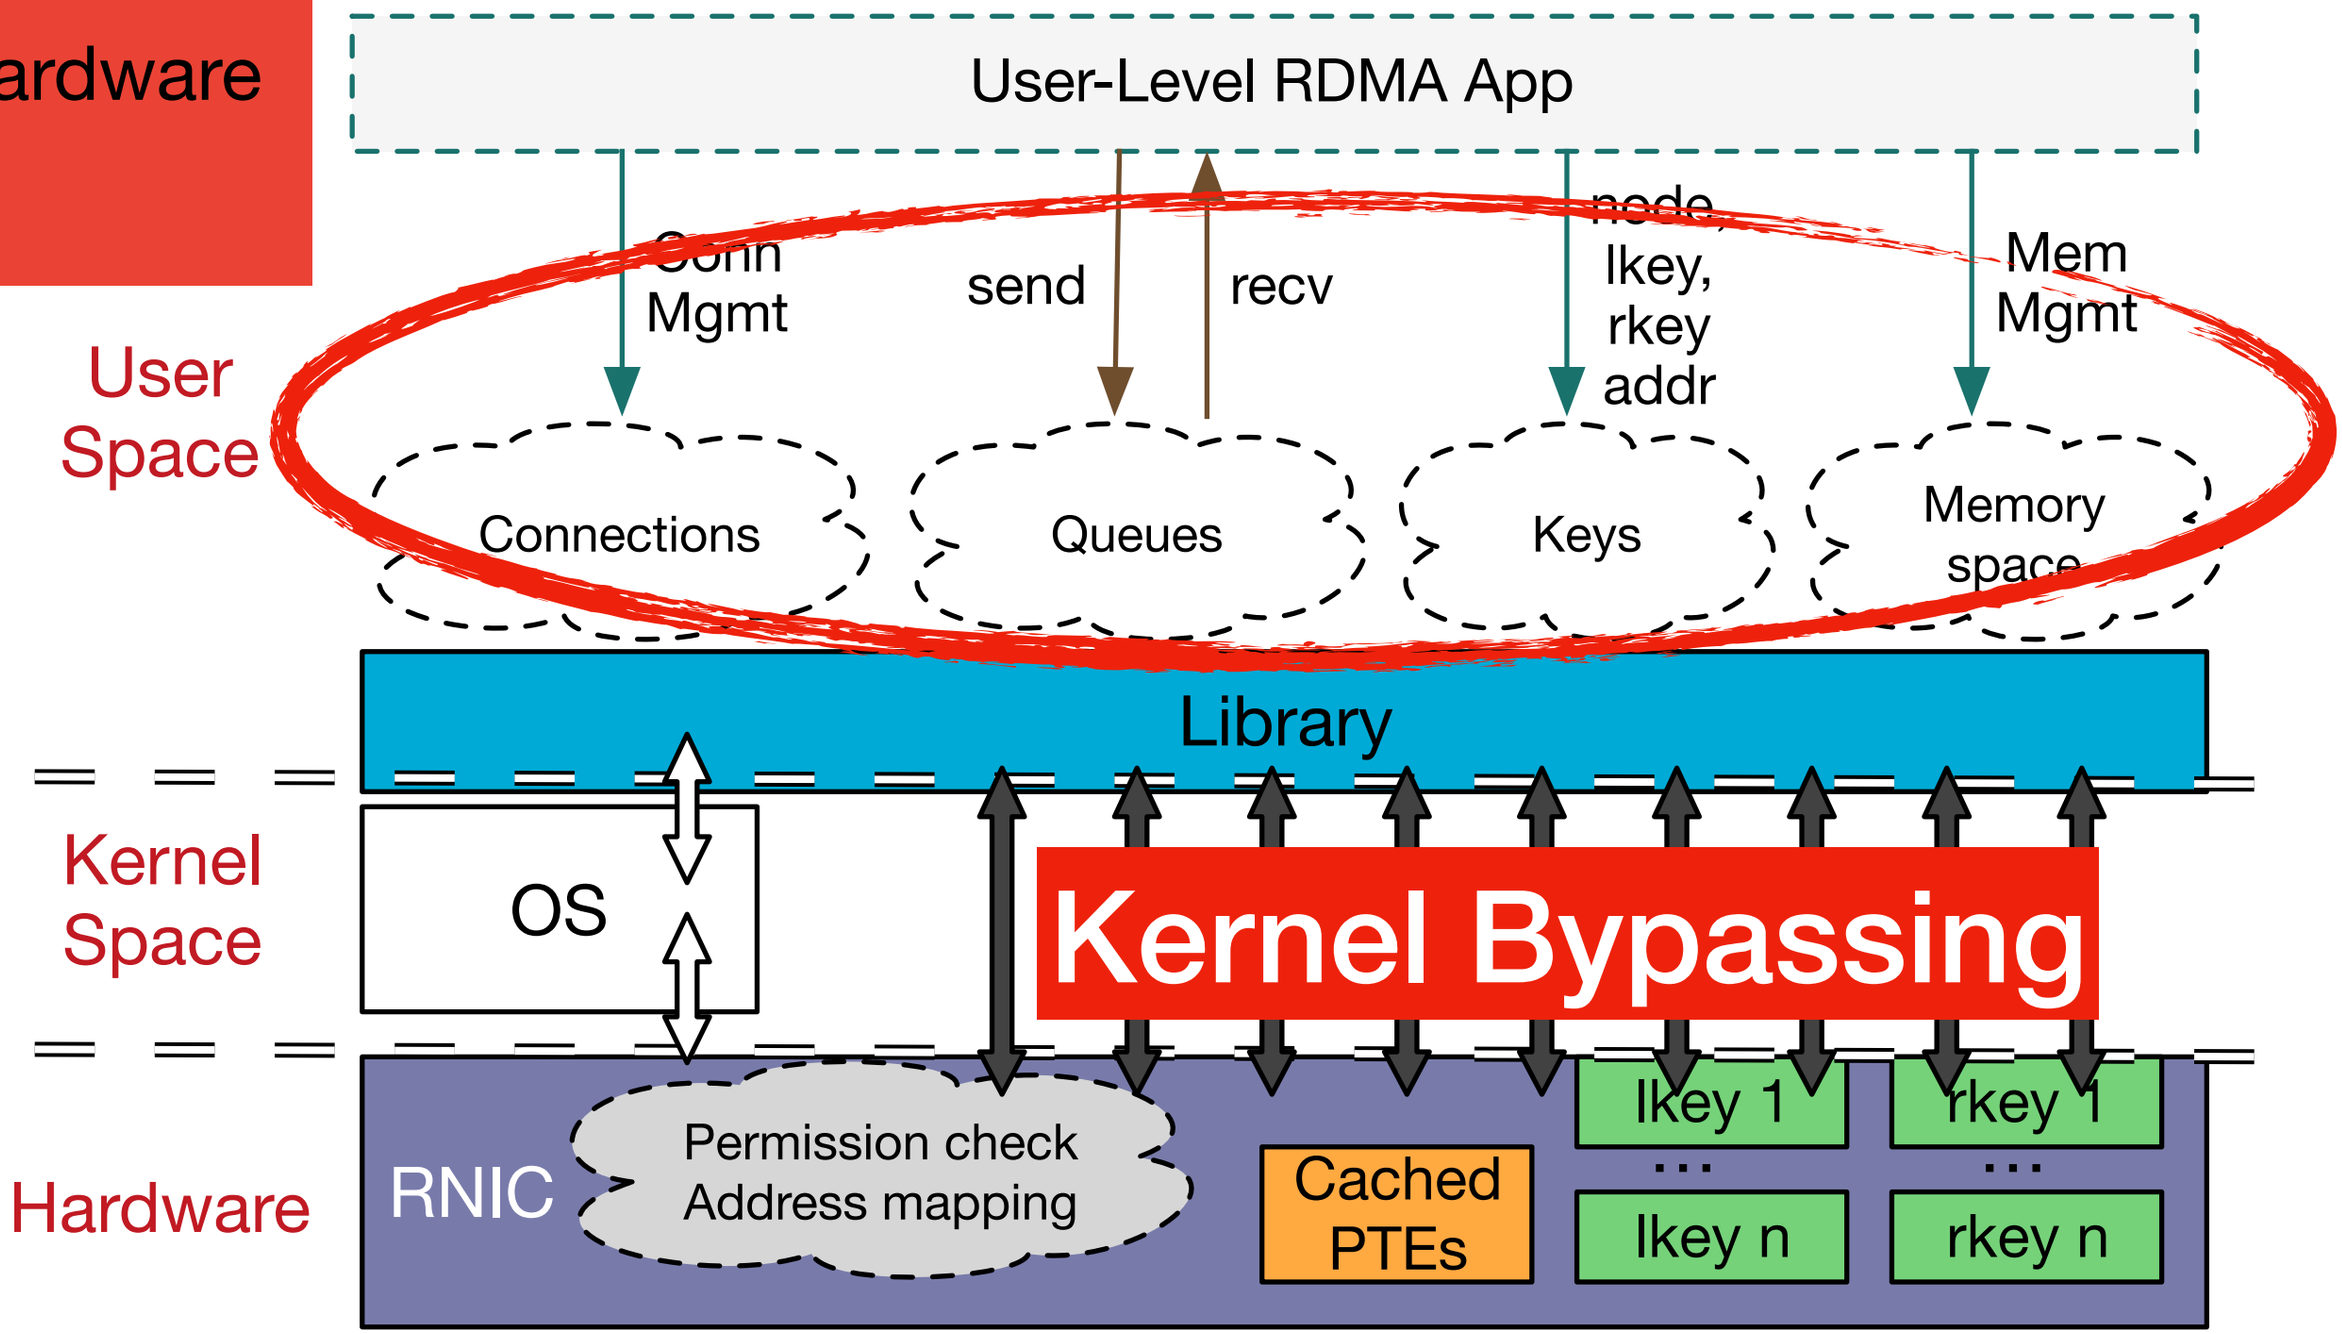

Userspace

Hardware

RDMA

User-Level RDMA App

User Space

Fat applications

No resource sharing

Library

Kernel Space

OS

Kernel Bypassing

Hardware

RNIC

Permission check
Address mapping

Cached PTEs

lkey 1

lkey n

rkey 1

rkey n

RDMA

Userspace

Hardware

User Space

Kernel Space

Hardware

User-Level RDMA App

node,

em

mt

Co

ry

e

Fat applications

No resource sharing

Library

OS

Kernel Bypassing

RNIC

Permission check
Address mapping

Cached PTEs

lkey 1

...

lkey n

rkey 1

...

rkey n

RDMA

Userspace

Hardware

User Space

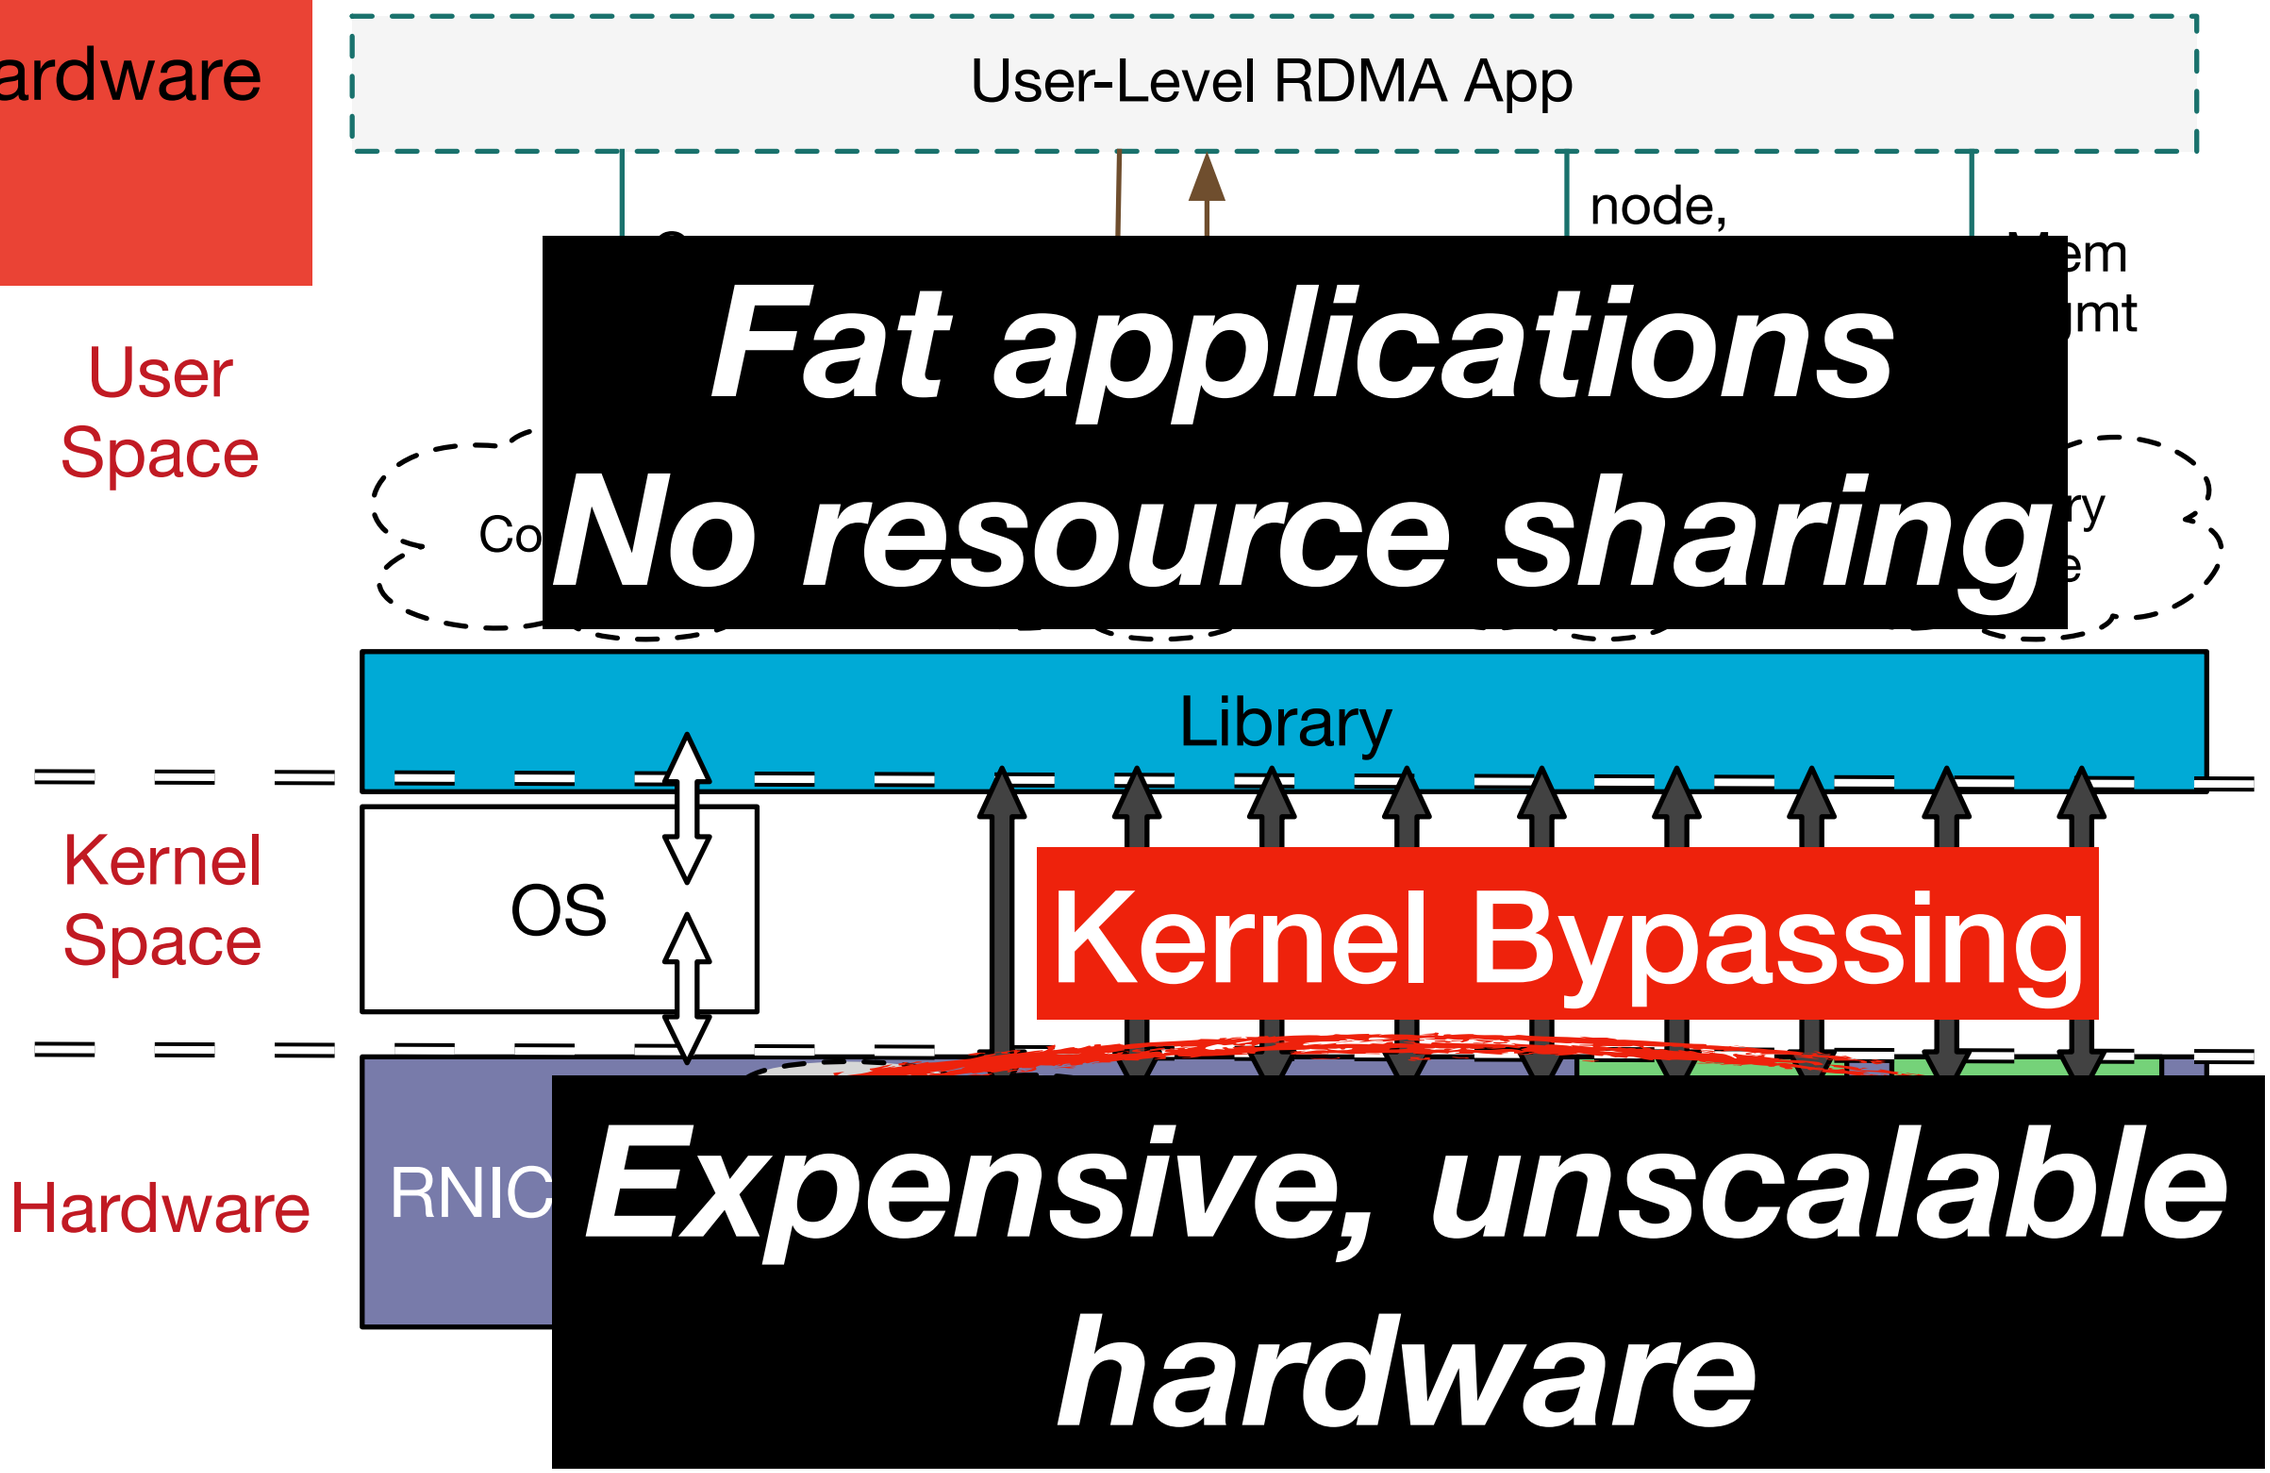

Fat applications

No resource sharing

Expensive, unscalable hardware

Are we removing too much from kernel?

Without Kernel

High-level abstraction

Resource sharing

Protection

Flexible QoS management

Without Kernel

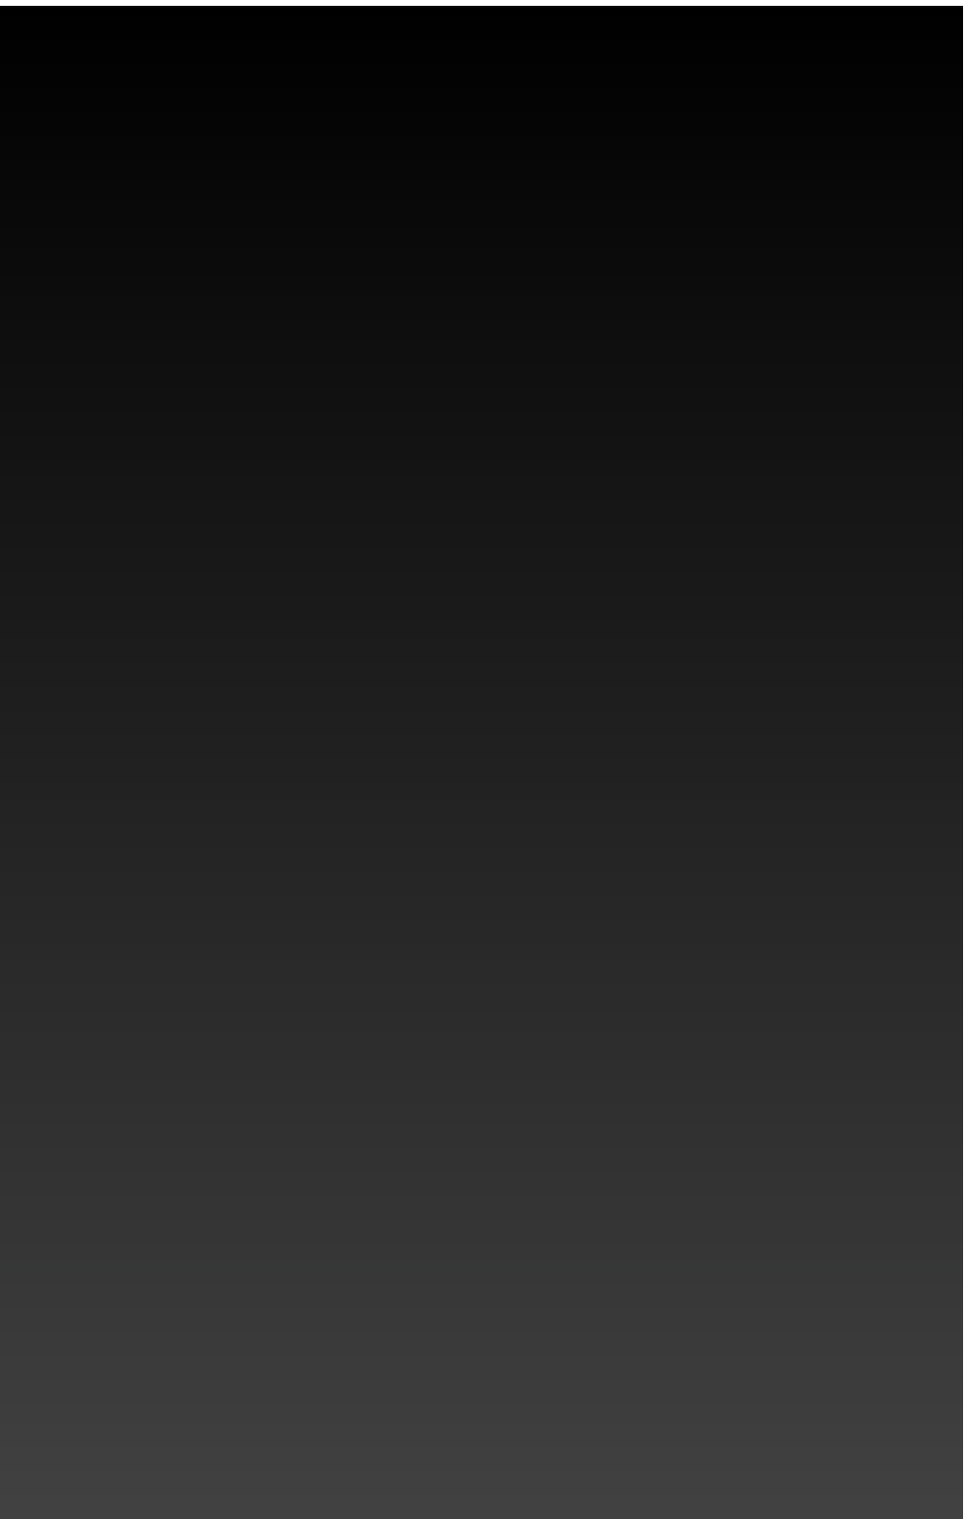

Resource sharing

Protection

Flexible QoS management

Without Kernel

Resource sharing

Flexible QoS management

Without Kernel

Flexible QoS management

Without Kernel

LITE - Kernel-Level Indirection for RDMA [SOSP'17]

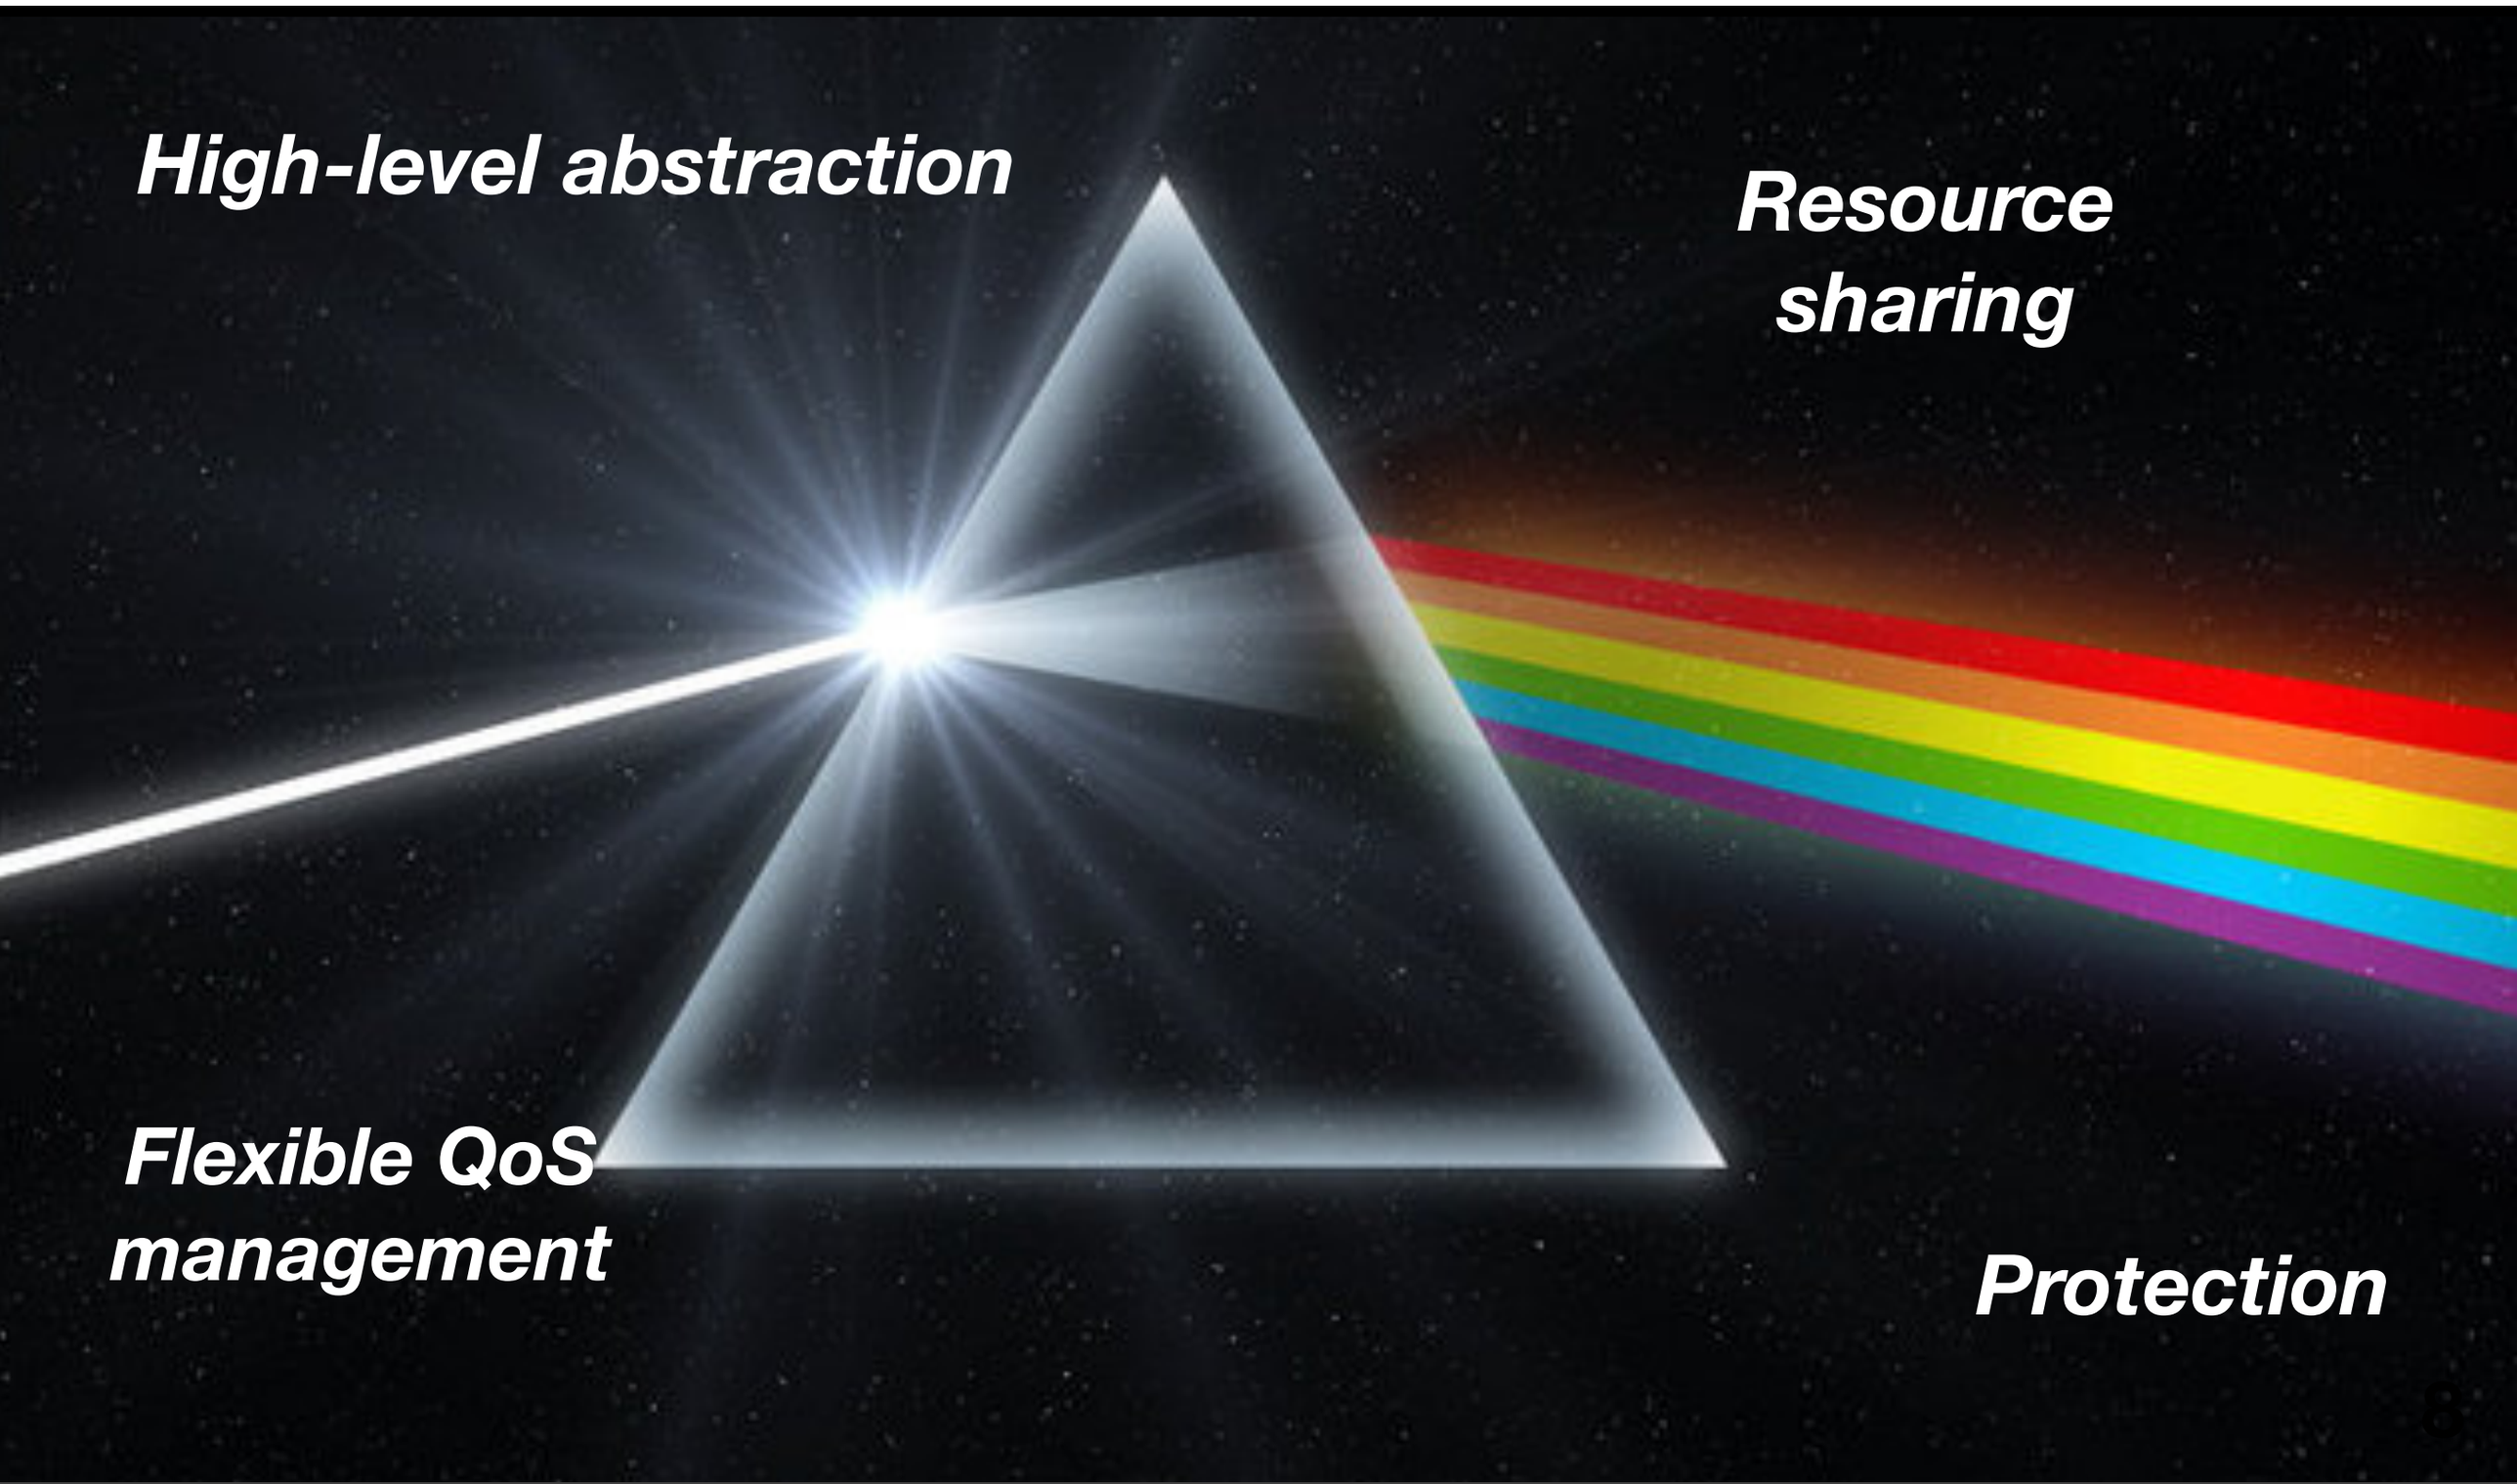

User-Level RDMA App

Conn Mgmt

send

recv

node, lkey, rkey addr

Mem Mgmt

User Space

Connections

Queues

Keys

Memory space

Library

Hardware

RNIC

Permission check
Address mapping

Cached PTEs

lkey 1
...
lkey n

rkey 1
...
rkey n

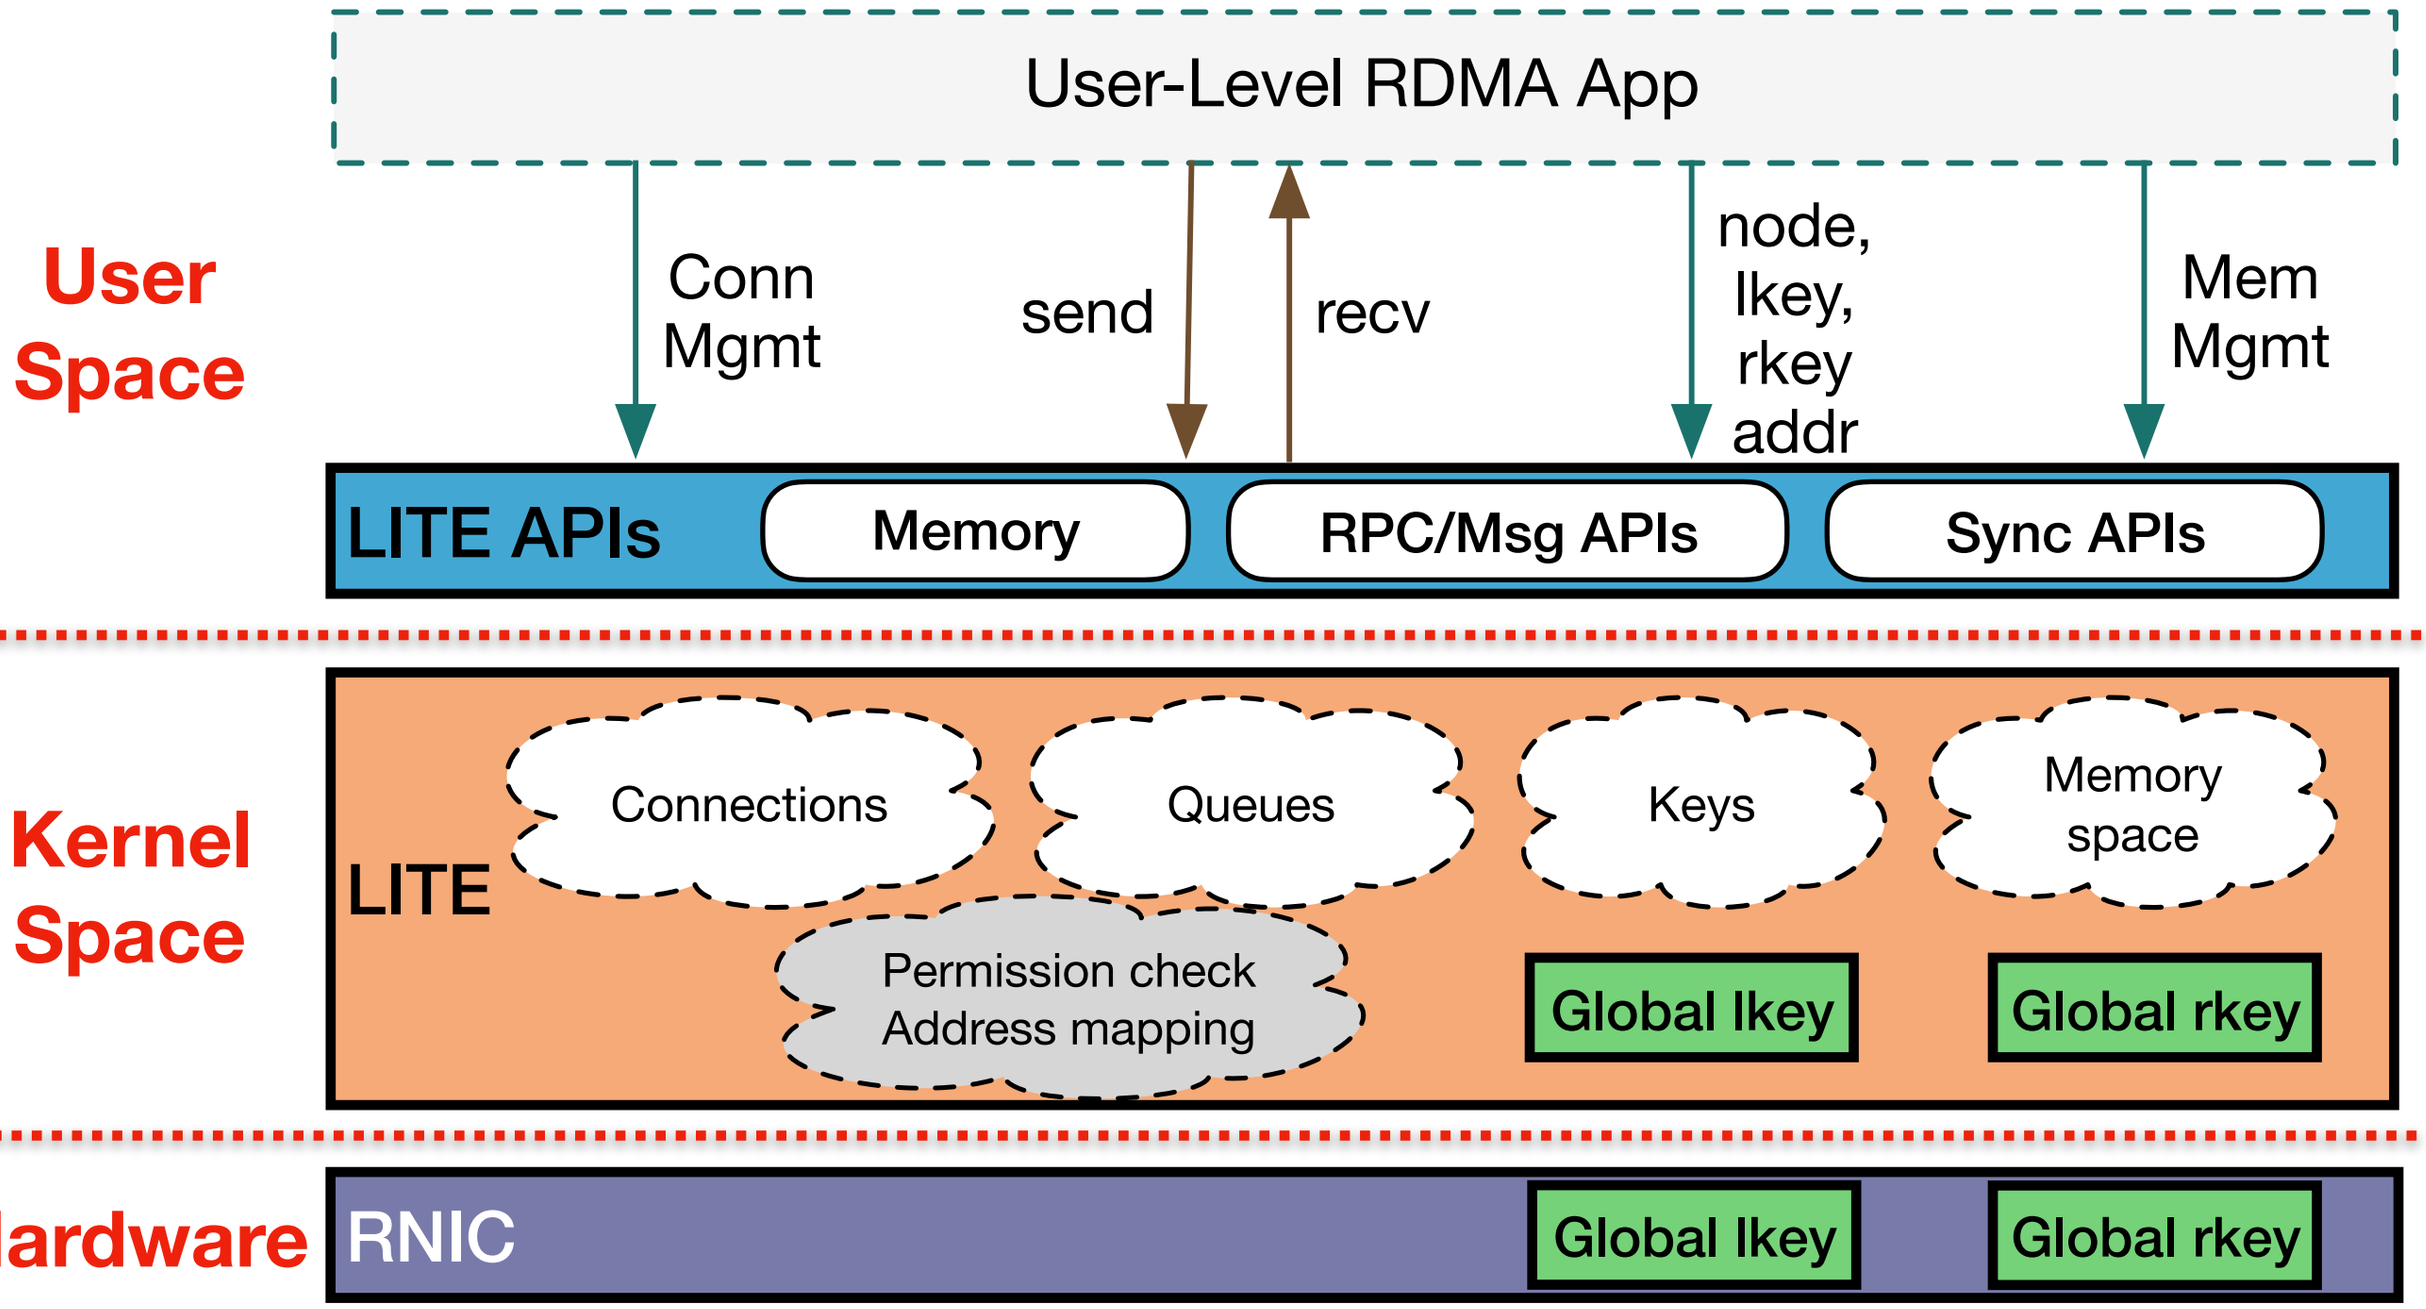
User-Level RDMA App
User Space
Conn Mgmt
send
recv
node, lkey, rkey addr
Mem Mgmt
LITE APIs
Memory
RPC/Msg APIs
Sync APIs
Kernel Space
LITE
Connections
Queues
Keys
Memory space
Permission check
Address mapping
Global lkey
Global rkey
Hardware
RNIC
Global lkey
Global rkey

Simpler applications

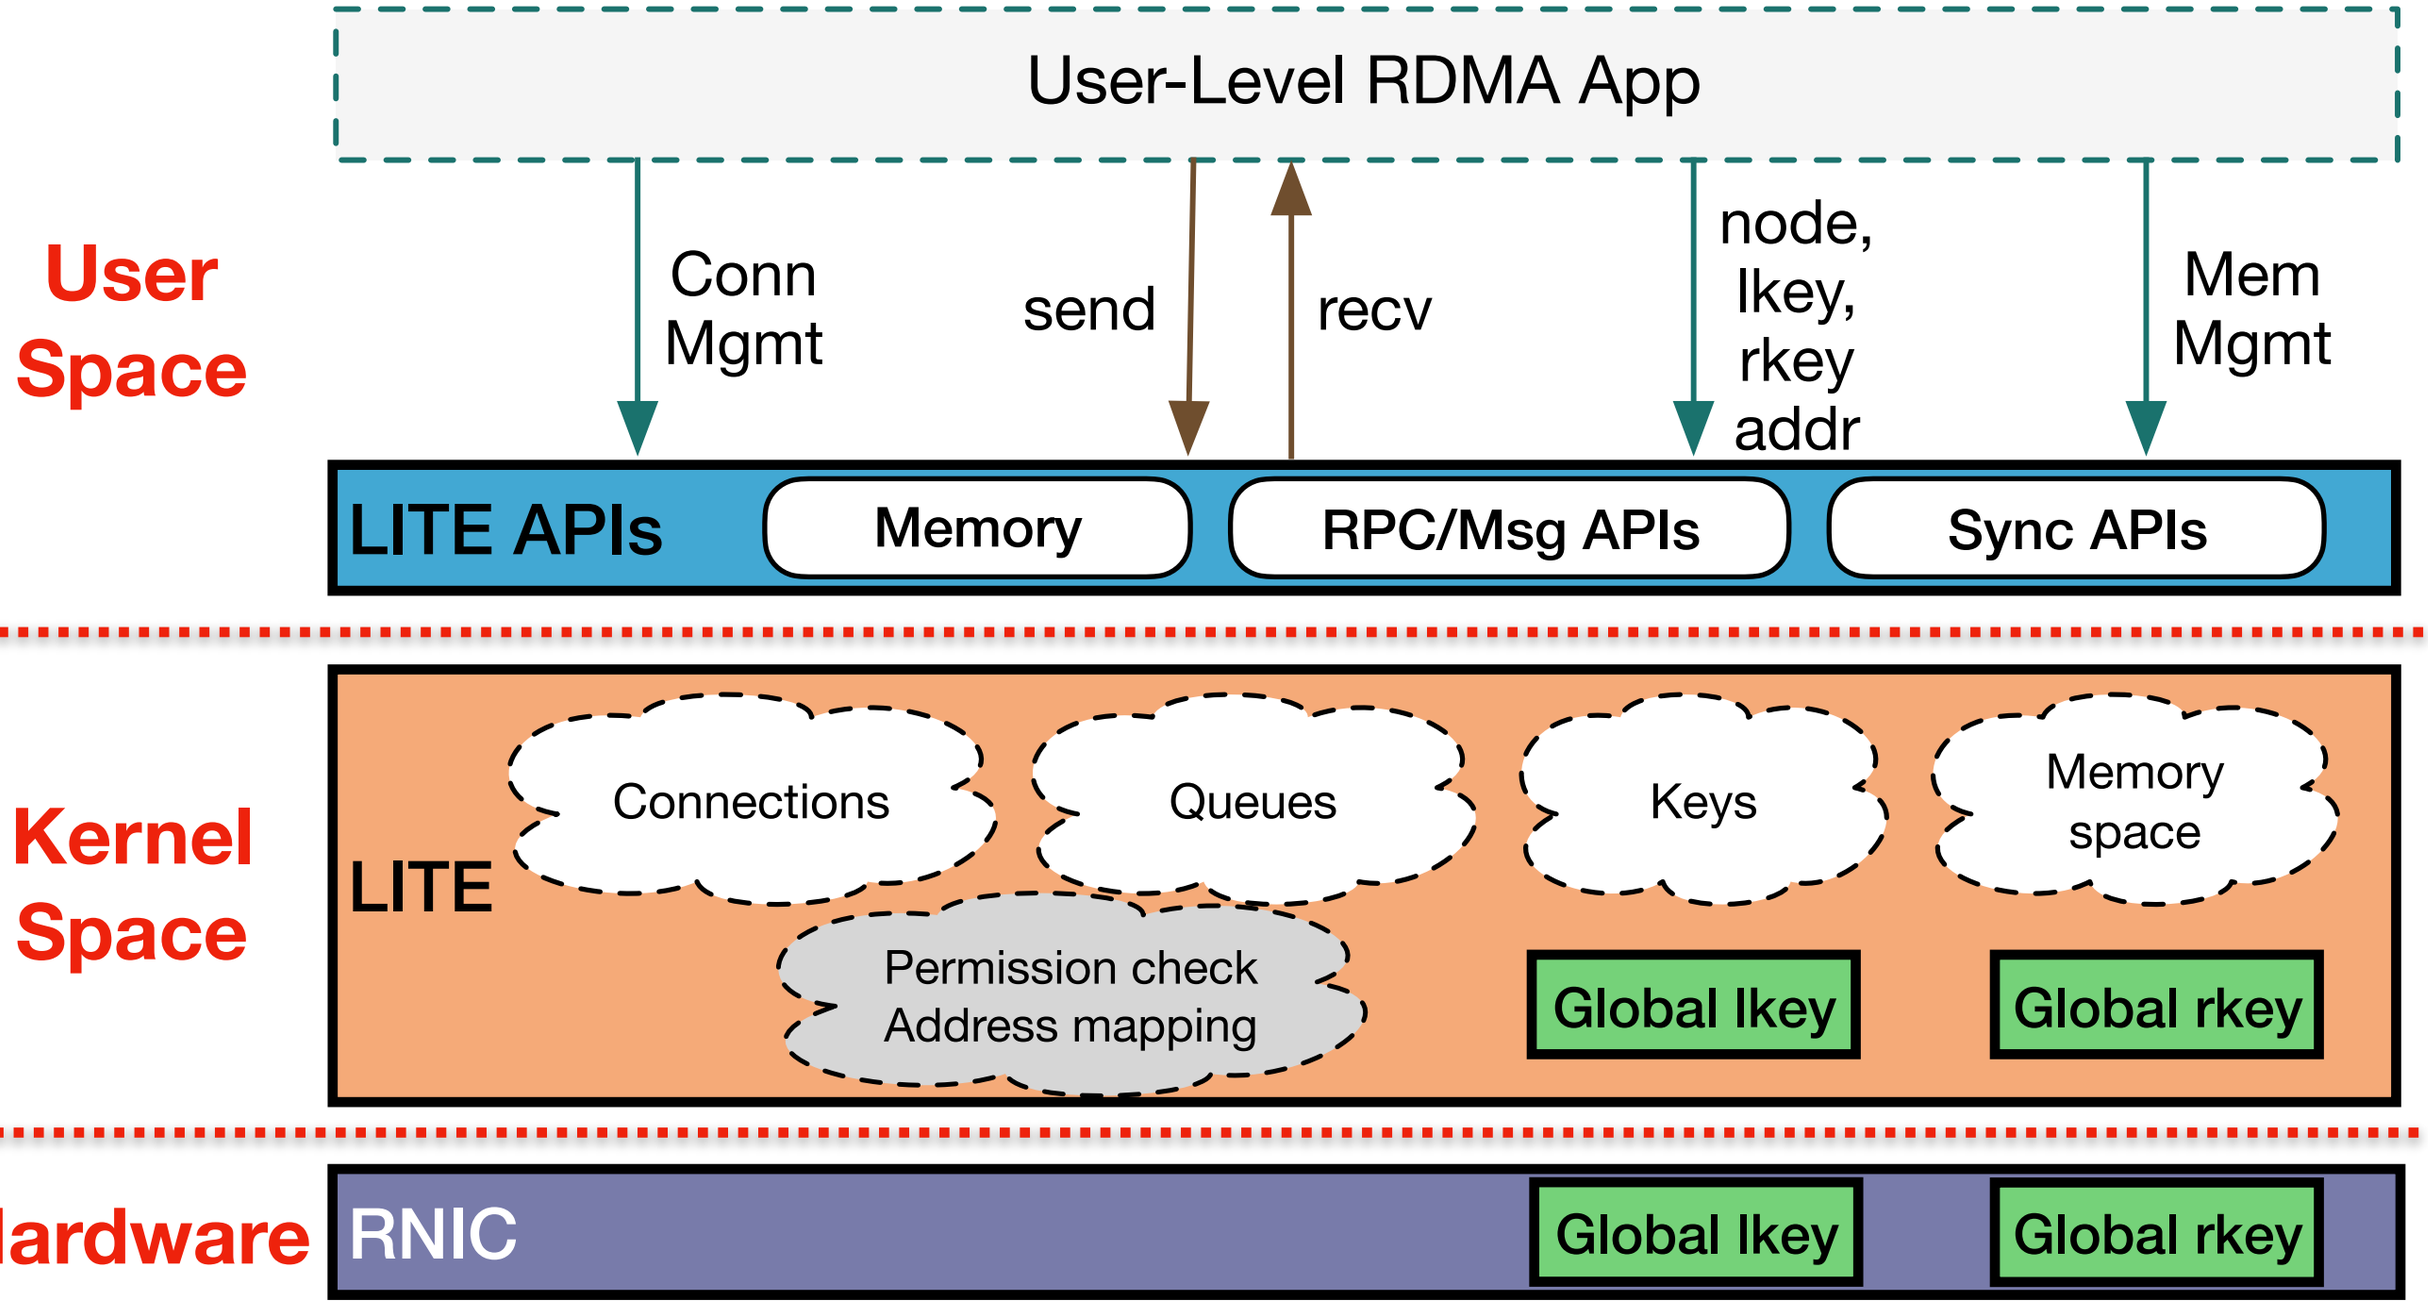

Simpler applications

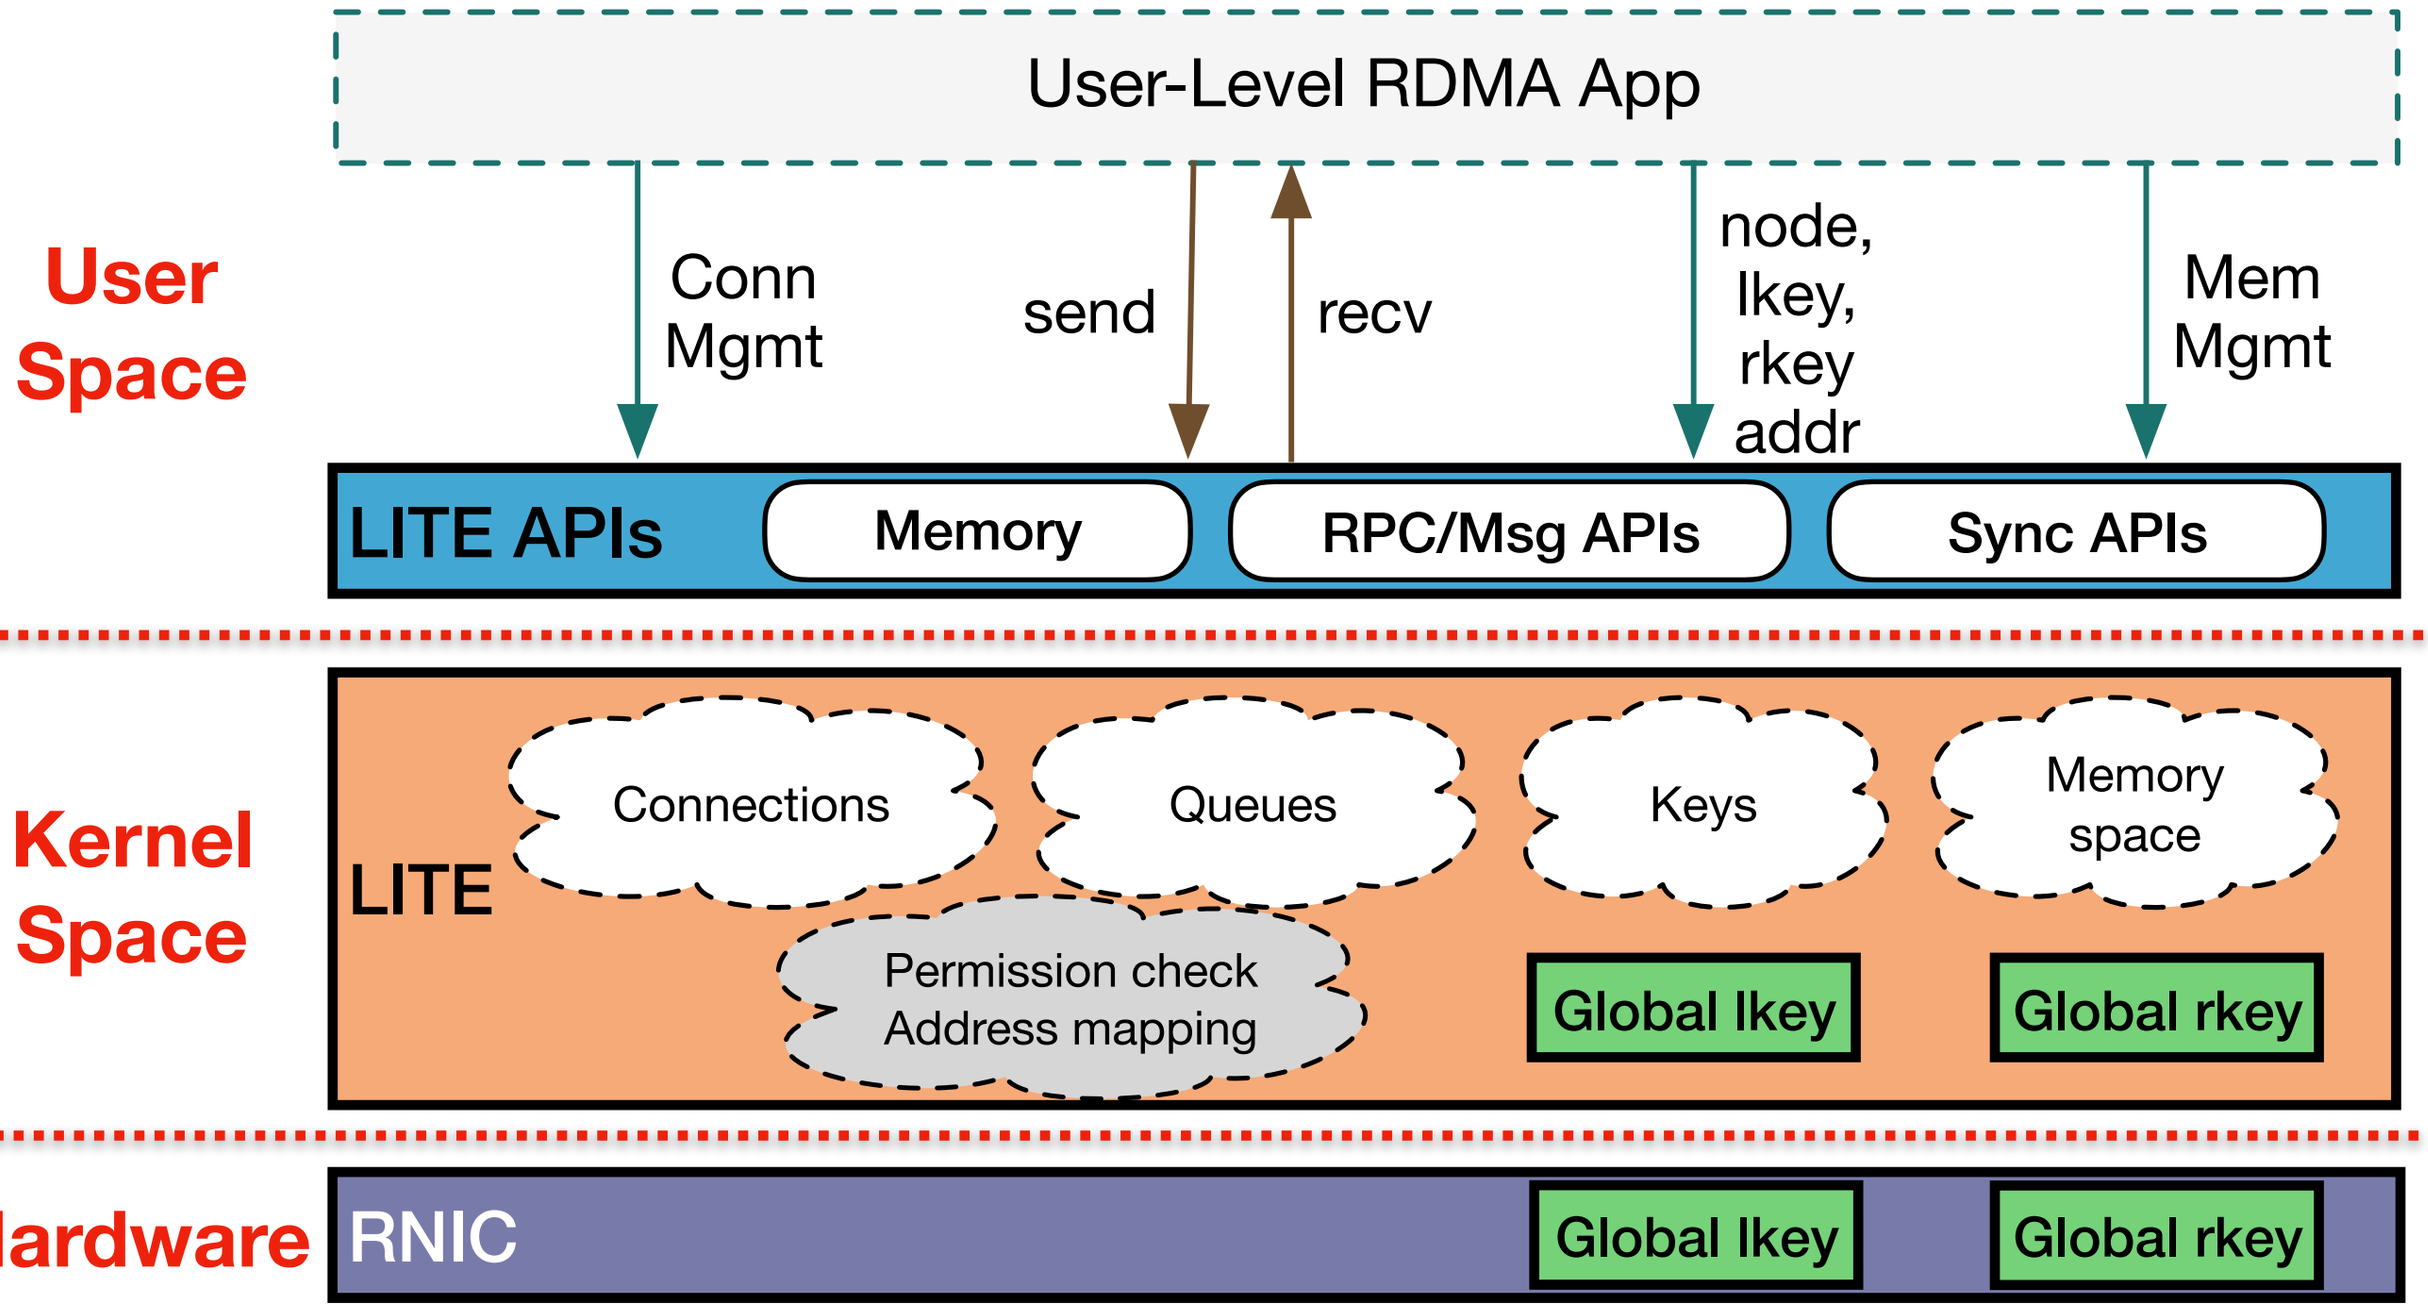

Cheaper hardware

Simpler applications

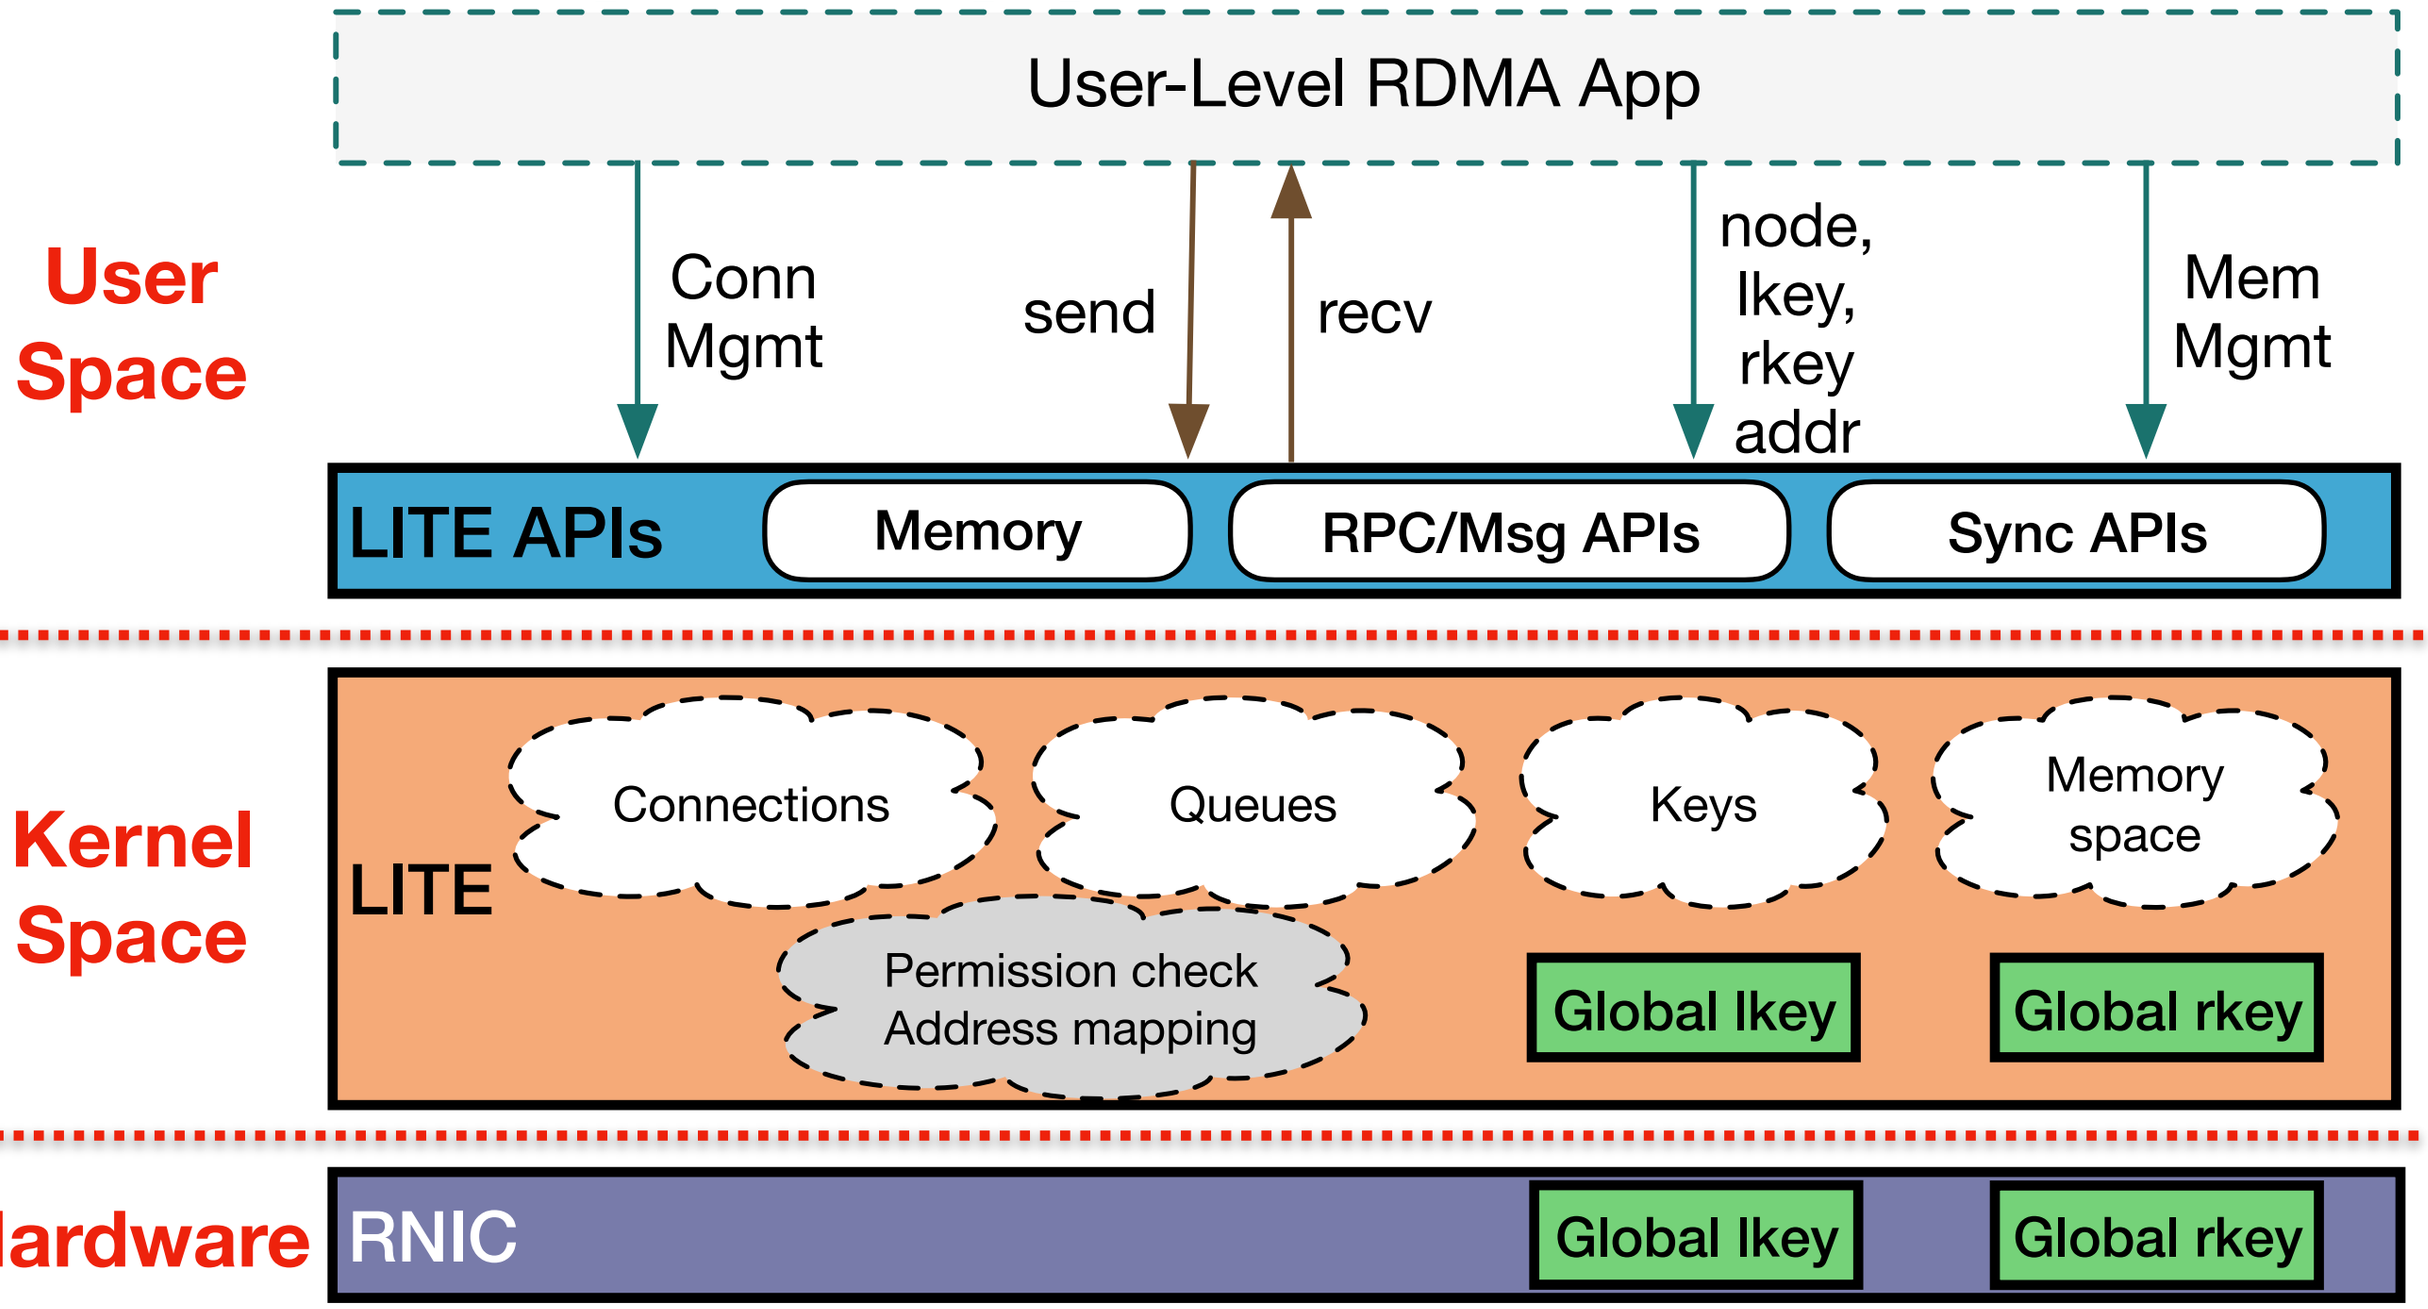

Cheaper hardware

Great Performance and scalability

My students, **before** going to their first conference:

Kernel programming is fun; we like doing great engineer work

Projects at WukLab

- ***LITE*** [sosp'17]: **15K** LOC, 80% in kernel
- ***Hotpot*** [SoCC'17]: **19K** LOC, all in kernel
- ***Lego*** OS [ongoing]: **170K** LOC already, all in kernel

My students, **after** going to their first conference:

Why are we writing so many kernel code when other students can get a paper with hundreds lines of user-level code?

Thank you
Questions?

Get LITE at: github.com/Wuklab/LITE

WukLab

wuklab.io

PURDUE UNIVERSITY